

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import math
import statsmodels.api as sm
```

```
In [3]: from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from xgboost import XGBRegressor
from sklearn.tree import plot_tree
from sklearn.inspection import PartialDependenceDisplay
from functools import reduce
```

```
In [4]: # Loading Dataset
df = pd.read_csv("train.csv")
```

```
In [5]: df.shape
```

```
Out[5]: (909604, 16)
```

```
In [6]: df.info()
```

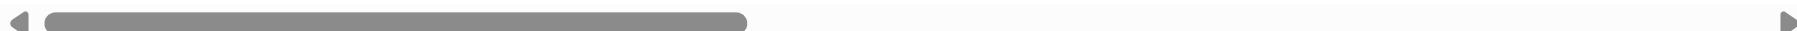
```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 909604 entries, 0 to 909603
Data columns (total 16 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   timestamp        909604 non-null   object  
 1   active_power_calculated_by_converter 909604 non-null   float64 
 2   active_power_raw      909604 non-null   float64 
 3   ambient_temperature    909604 non-null   float64 
 4   generator_speed       909604 non-null   float64 
 5   generator_winding_temp_max 909604 non-null   float64 
 6   grid_power10min_average 909604 non-null   float64 
 7   nc1_inside_temp       909604 non-null   float64 
 8   nacelle_temp          909604 non-null   float64 
 9   reactice_power_calculated_by_converter 909604 non-null   float64 
 10  reactive_power        909604 non-null   float64 
 11  wind_direction_raw     909604 non-null   float64 
 12  wind_speed_raw         909604 non-null   float64 
 13  wind_speed_turbulence   909604 non-null   float64 
 14  turbine_id            909604 non-null   object  
 15  Target                909604 non-null   float64 

dtypes: float64(14), object(2)
memory usage: 111.0+ MB
```

In [7]: `df.head()`

Out[7]:

	timestamp	active_power_calculated_by_converter	active_power_raw	ambient_temperature	generator_speed	generator_winding_temp_max
0	2021-02-19 20:18:00	816.636759	834.917206	31.694380	1159.616602	65.95421
1	2021-04-27 04:55:00	419.107829	421.050873	12.894948	928.747996	59.57131
2	2021-01-25 06:26:00	1303.530558	1337.566142	16.648388	1201.219775	61.27045
3	2021-10-30 03:47:00	61.494872	53.481008	28.388141	769.806122	40.67434
4	2021-03-15 00:39:00	593.514364	611.659108	31.519527	1046.916768	64.34176

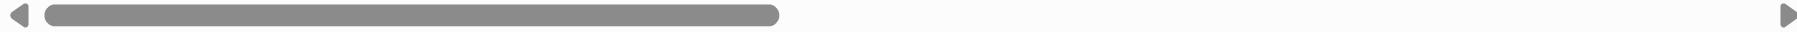


In [8]:

```
df.describe()
```

Out[8]:

	active_power_calculated_by_converter	active_power_raw	ambient_temperature	generator_speed	generator_winding_temp_max	grid_
count	909604.000000	909604.000000	909604.000000	909604.000000	909604.000000	
mean	525.886061	537.639468	25.573921	931.130882	60.730112	
std	474.619576	482.394907	8.025649	193.674384	9.721582	
min	0.000000	0.000008	5.616541	0.000000	25.636366	
25%	149.620898	155.813244	17.093680	770.538574	54.916767	
50%	379.989929	383.985926	27.919035	884.068502	60.034490	
75%	781.357475	800.261108	32.171696	1123.335856	65.737187	
max	1999.999858	1999.984456	48.089017	1267.140625	129.846405	



In [9]:

```
# Data Cleaning
```

```
In [10]: # Drop Turbine_id  
df = df.drop('turbine_id', axis = 1)
```

```
In [11]: # Convert timestamp to datetime  
df['timestamp'] = pd.to_datetime(df['timestamp'])
```

```
In [12]: # Extract time features  
df['hour'] = df['timestamp'].dt.hour  
df['day_of_week'] = df['timestamp'].dt.dayofweek  
df['month'] = df['timestamp'].dt.month  
df['season'] = df['month']%12 // 3 + 1
```

```
In [13]: df = df.drop('timestamp', axis = 1)
```

```
In [14]: df.head()
```

```
Out[14]: active_power_calculated_by_converter active_power_raw ambient_temperature generator_speed generator_winding_temp_max grid_pow  
0 816.636759 834.917206 31.694380 1159.616602 65.954214  
1 419.107829 421.050873 12.894948 928.747996 59.571319  
2 1303.530558 1337.566142 16.648388 1201.219775 61.270498  
3 61.494872 53.481008 28.388141 769.806122 40.674348  
4 593.514364 611.659108 31.519527 1046.916768 64.341763
```

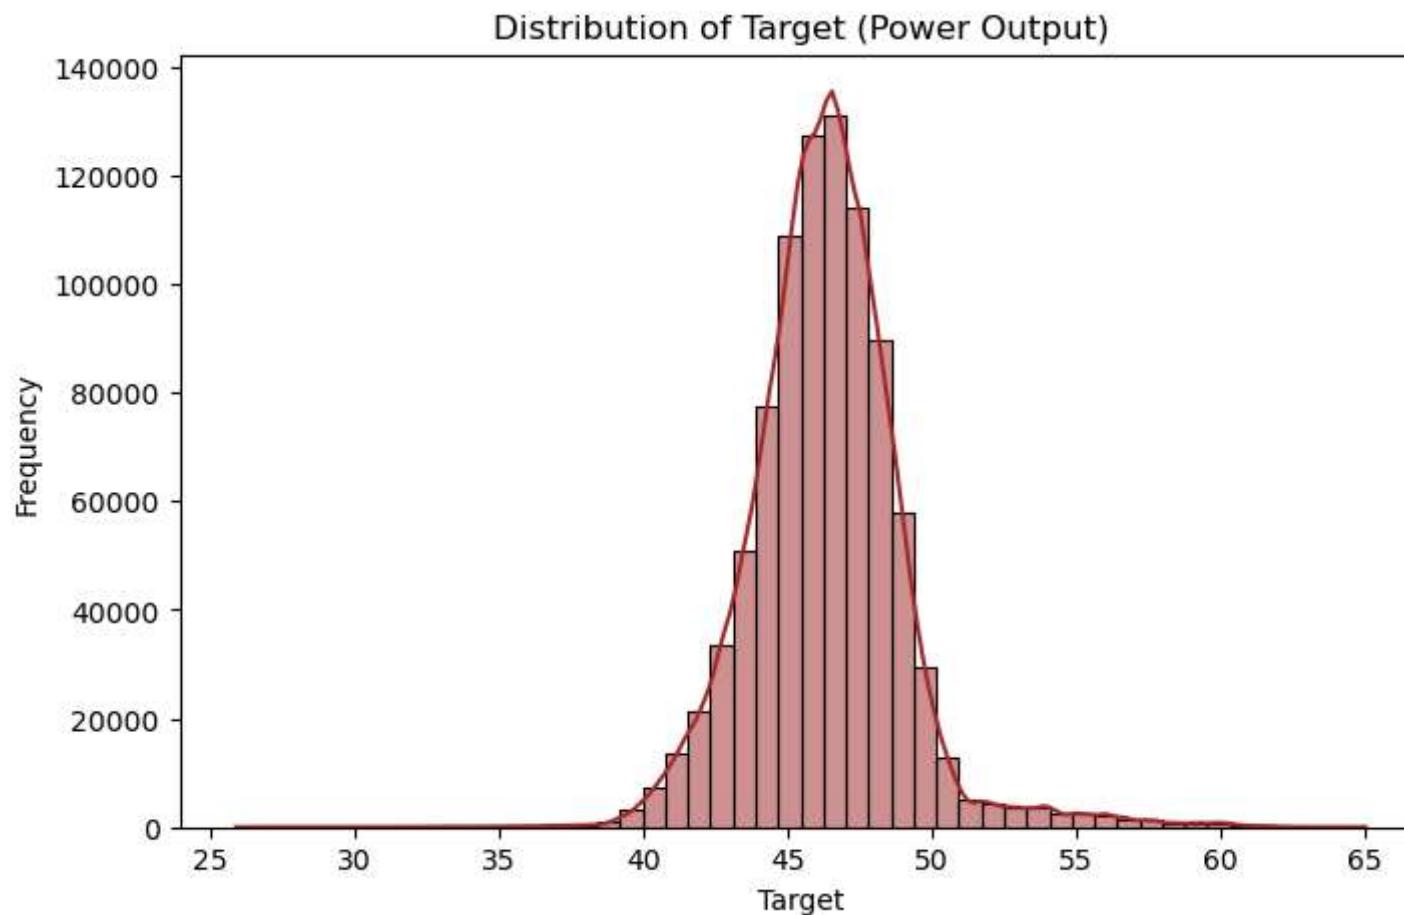


```
In [15]: # Checking Null or missing values in each column  
df.isnull().sum()
```

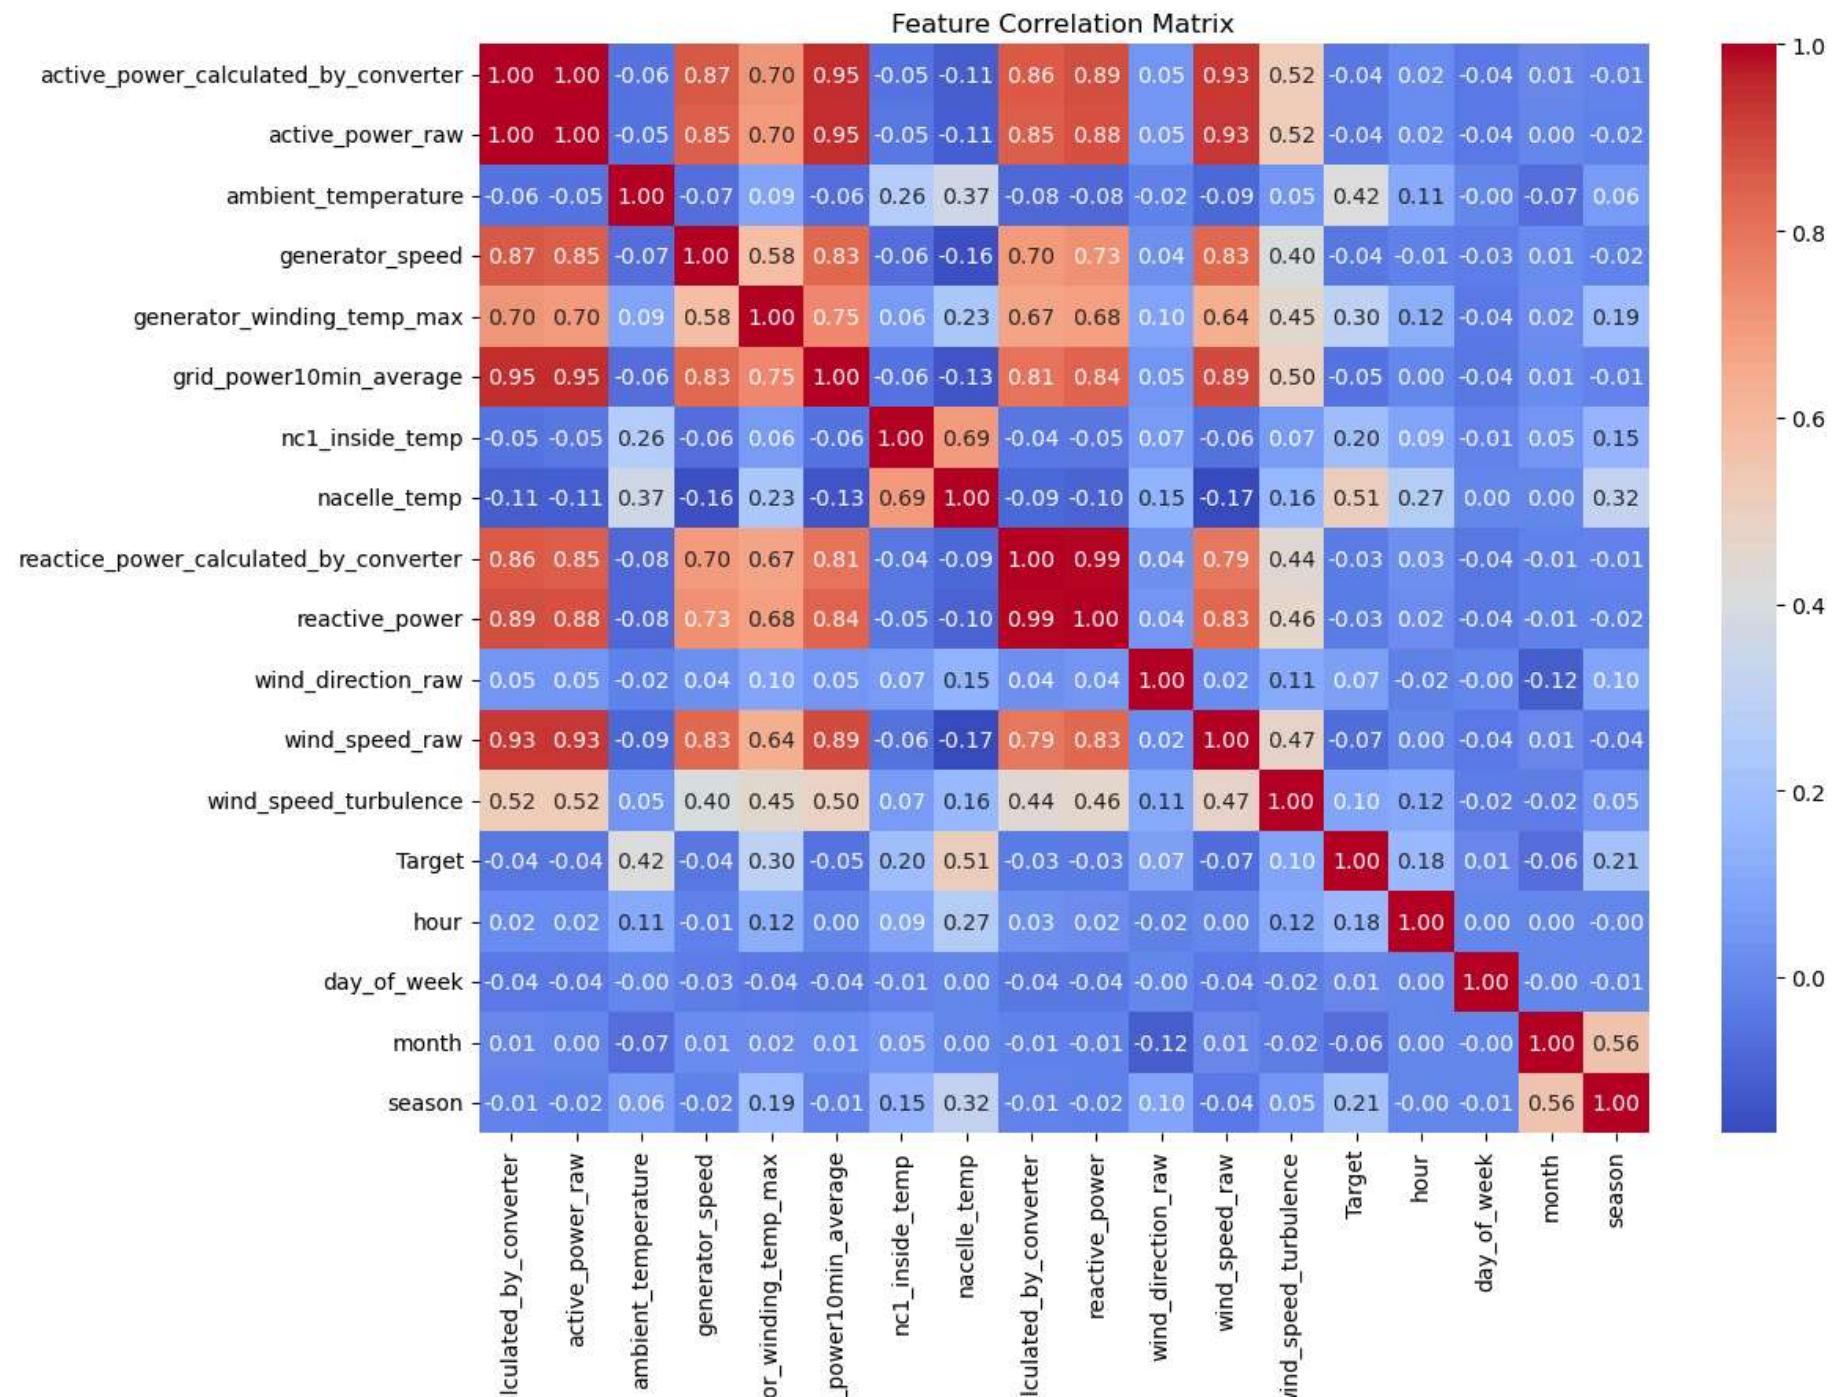
```
Out[15]: active_power_calculated_by_converter      0
          active_power_raw                      0
          ambient_temperature                  0
          generator_speed                   0
          generator_winding_temp_max        0
          grid_power10min_average         0
          nc1_inside_temp                  0
          nacelle_temp                     0
          reactice_power_calculated_by_converter 0
          reactive_power                   0
          wind_direction_raw              0
          wind_speed_raw                  0
          wind_speed_turbulence           0
          Target                          0
          hour                            0
          day_of_week                     0
          month                           0
          season                          0
          dtype: int64
```

```
In [16]: # Exploratory Data Analysis ( EDA )
```

```
In [17]: # Target Distribution
plt.figure(figsize = (8, 5))
sns.histplot(df['Target'], bins = 50, kde = True, color = 'brown')
plt.title('Distribution of Target (Power Output)')
plt.xlabel('Target')
plt.ylabel('Frequency')
plt.show()
```



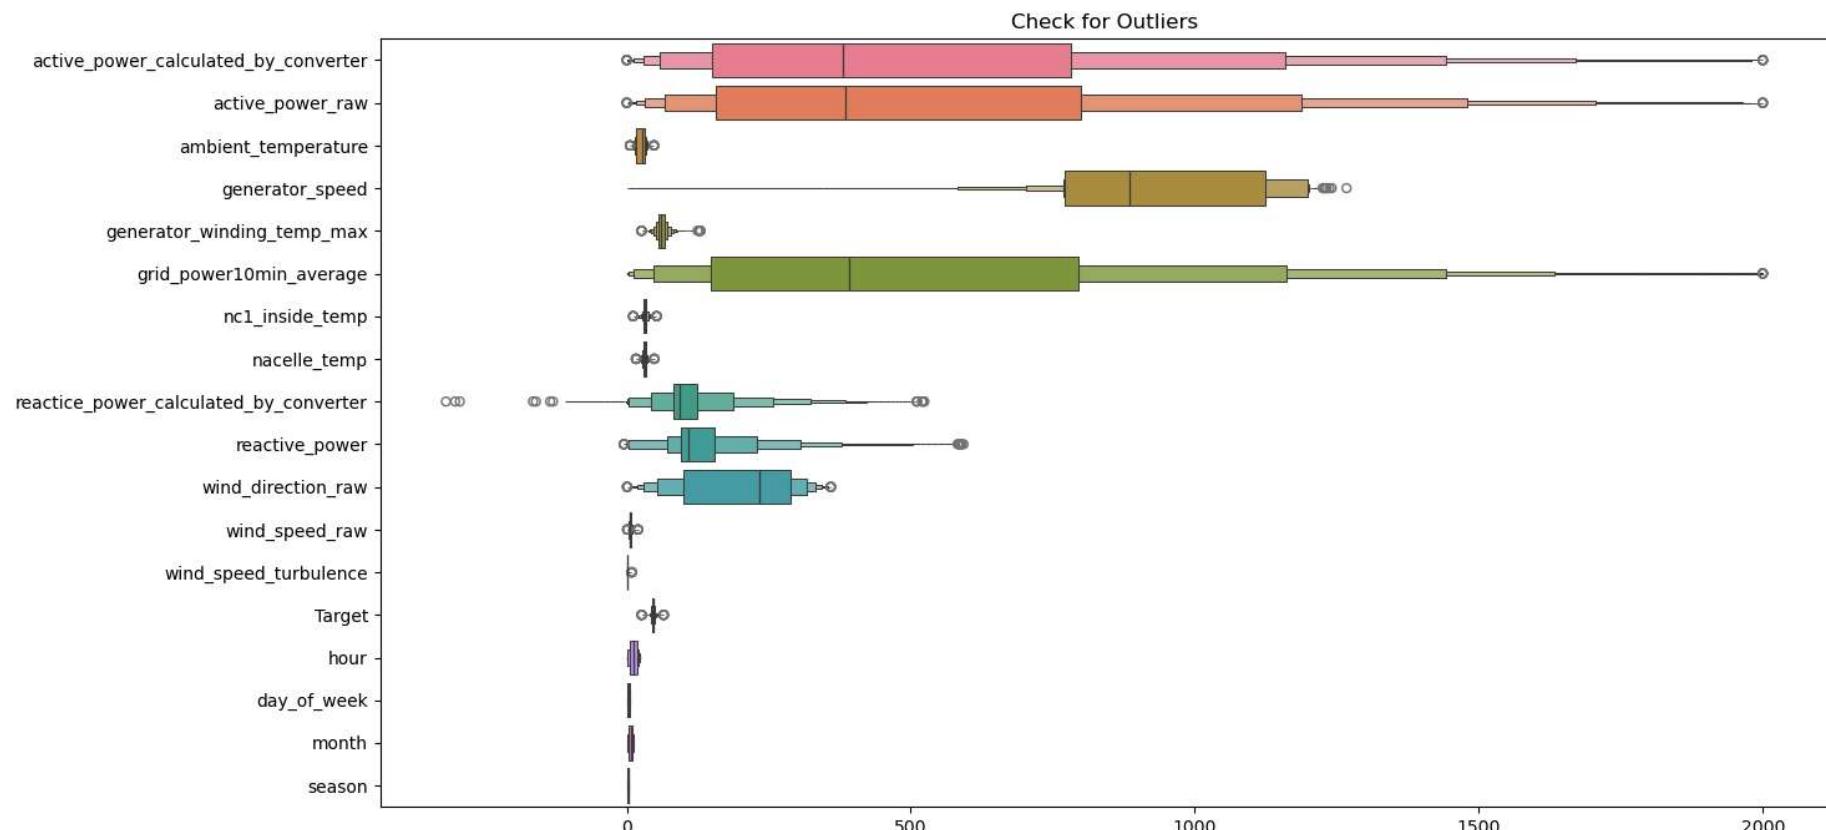
```
In [18]: # Feature Correlation Matrix
corr = df.corr()
plt.figure(figsize = (12, 9))
sns.heatmap(corr, annot = True, fmt = ".2f", cmap = 'coolwarm')
plt.title('Feature Correlation Matrix')
plt.show()
```



w

In [19]:

```
# Checking For Outliers
numeric_features = df.select_dtypes(include = np.number).columns.tolist()
plt.figure(figsize = (15, 8))
sns.boxenplot(data = df[numeric_features], orient = 'h')
plt.title('Check for Outliers')
# plt.savefig("outlier.png", dpi=300, bbox_inches="tight")
plt.show()
```



```
In [20]: # Select all numeric features except Target
```

```
numeric_features = df.select_dtypes(include = np.number).columns.tolist()
numeric_features.remove('Target')
print("Numeric features:", numeric_features)
```

```
Numeric features: ['active_power_calculated_by_converter', 'active_power_raw', 'ambient_temperature', 'generator_speed', 'generator_winding_temp_max', 'grid_power10min_average', 'nc1_inside_temp', 'nacelle_temp', 'reactice_power_calculated_by_converter', 'reactive_power', 'wind_direction_raw', 'wind_speed_raw', 'wind_speed_turbulence', 'hour', 'day_of_week', 'month', 'season']
```

```
In [21]: # Blancing the extreme outliers and scaling numeric features
```

```
In [22]: # Removing extreme outliers
```

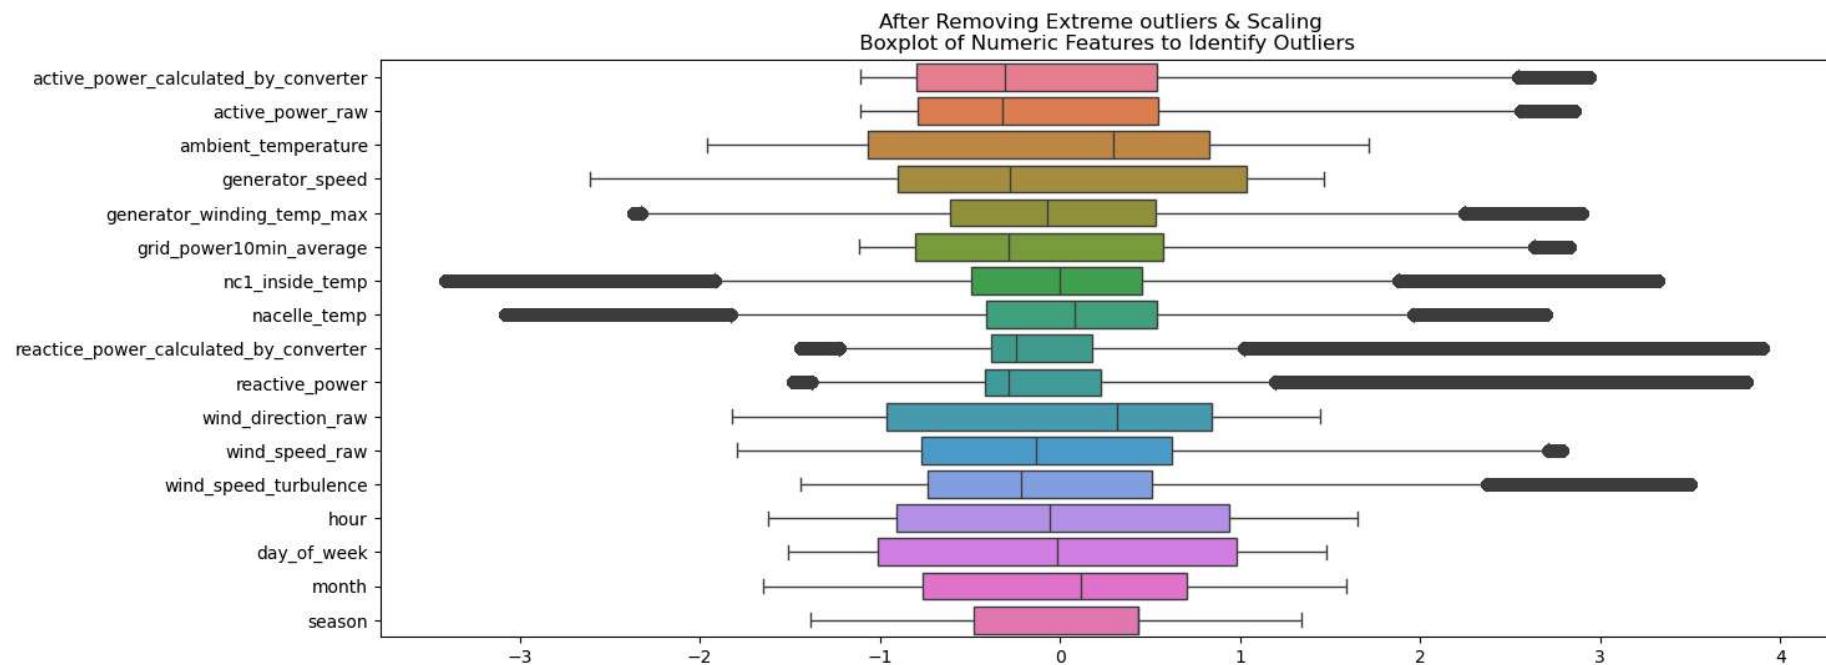
```
for col in numeric_features:
    lower = df[col].quantile(0.01)
    upper = df[col].quantile(0.99)
    df[col] = df[col].clip(lower, upper)
```

```
In [23]: # Scaling Numeric features
```

```
scaler = StandardScaler()
df[numeric_features] = scaler.fit_transform(df[numeric_features])
```

```
In [24]: # Checking Outliers Again
```

```
plt.figure(figsize=(15,6))
sns.boxplot(data=df[numeric_features], orient='h')
plt.title('After Removing Extreme outliers & Scaling \n Boxplot of Numeric Features to Identify Outliers')
plt.show()
```



In [25]: # Features vs Target

```
# Select only numeric features (excluding Target)
numeric_features = df.select_dtypes(include=np.number).columns.tolist()
numeric_features.remove('Target')

# Set up grid Layout (2 columns)
n_cols = 2
n_rows = math.ceil(len(numeric_features) / n_cols)

plt.style.use('seaborn-v0_8-whitegrid')
fig, axes = plt.subplots(n_rows, n_cols, figsize=(12, n_rows * 3), constrained_layout = True)
axes = axes.flatten() # flatten for easy indexing

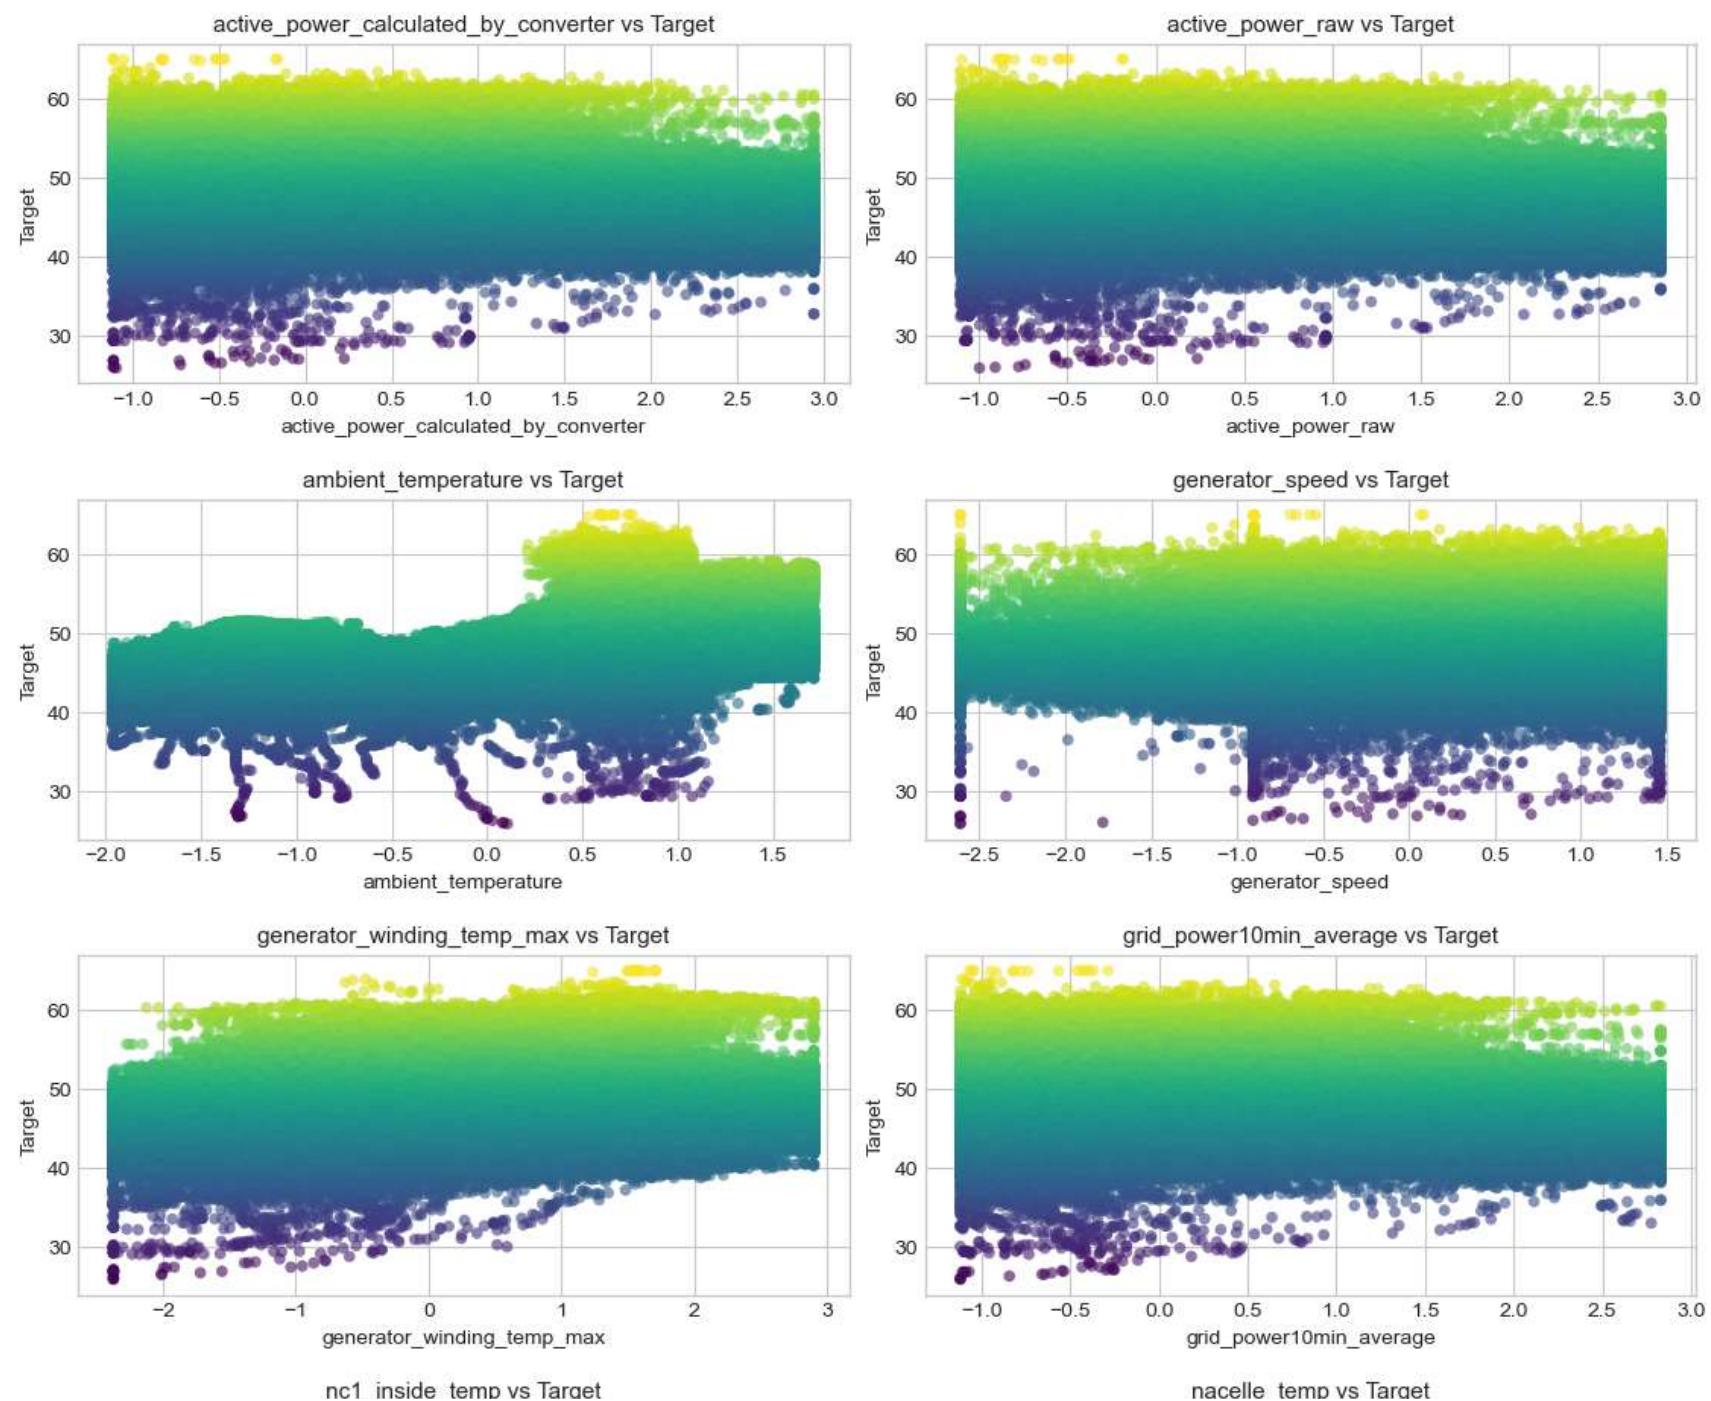
for i, col in enumerate(numeric_features):
    sc = axes[i].scatter(df[col], df['Target'],
                         c=df['Target'],
                         cmap='viridis',
                         alpha=0.6,
                         s=30,
```

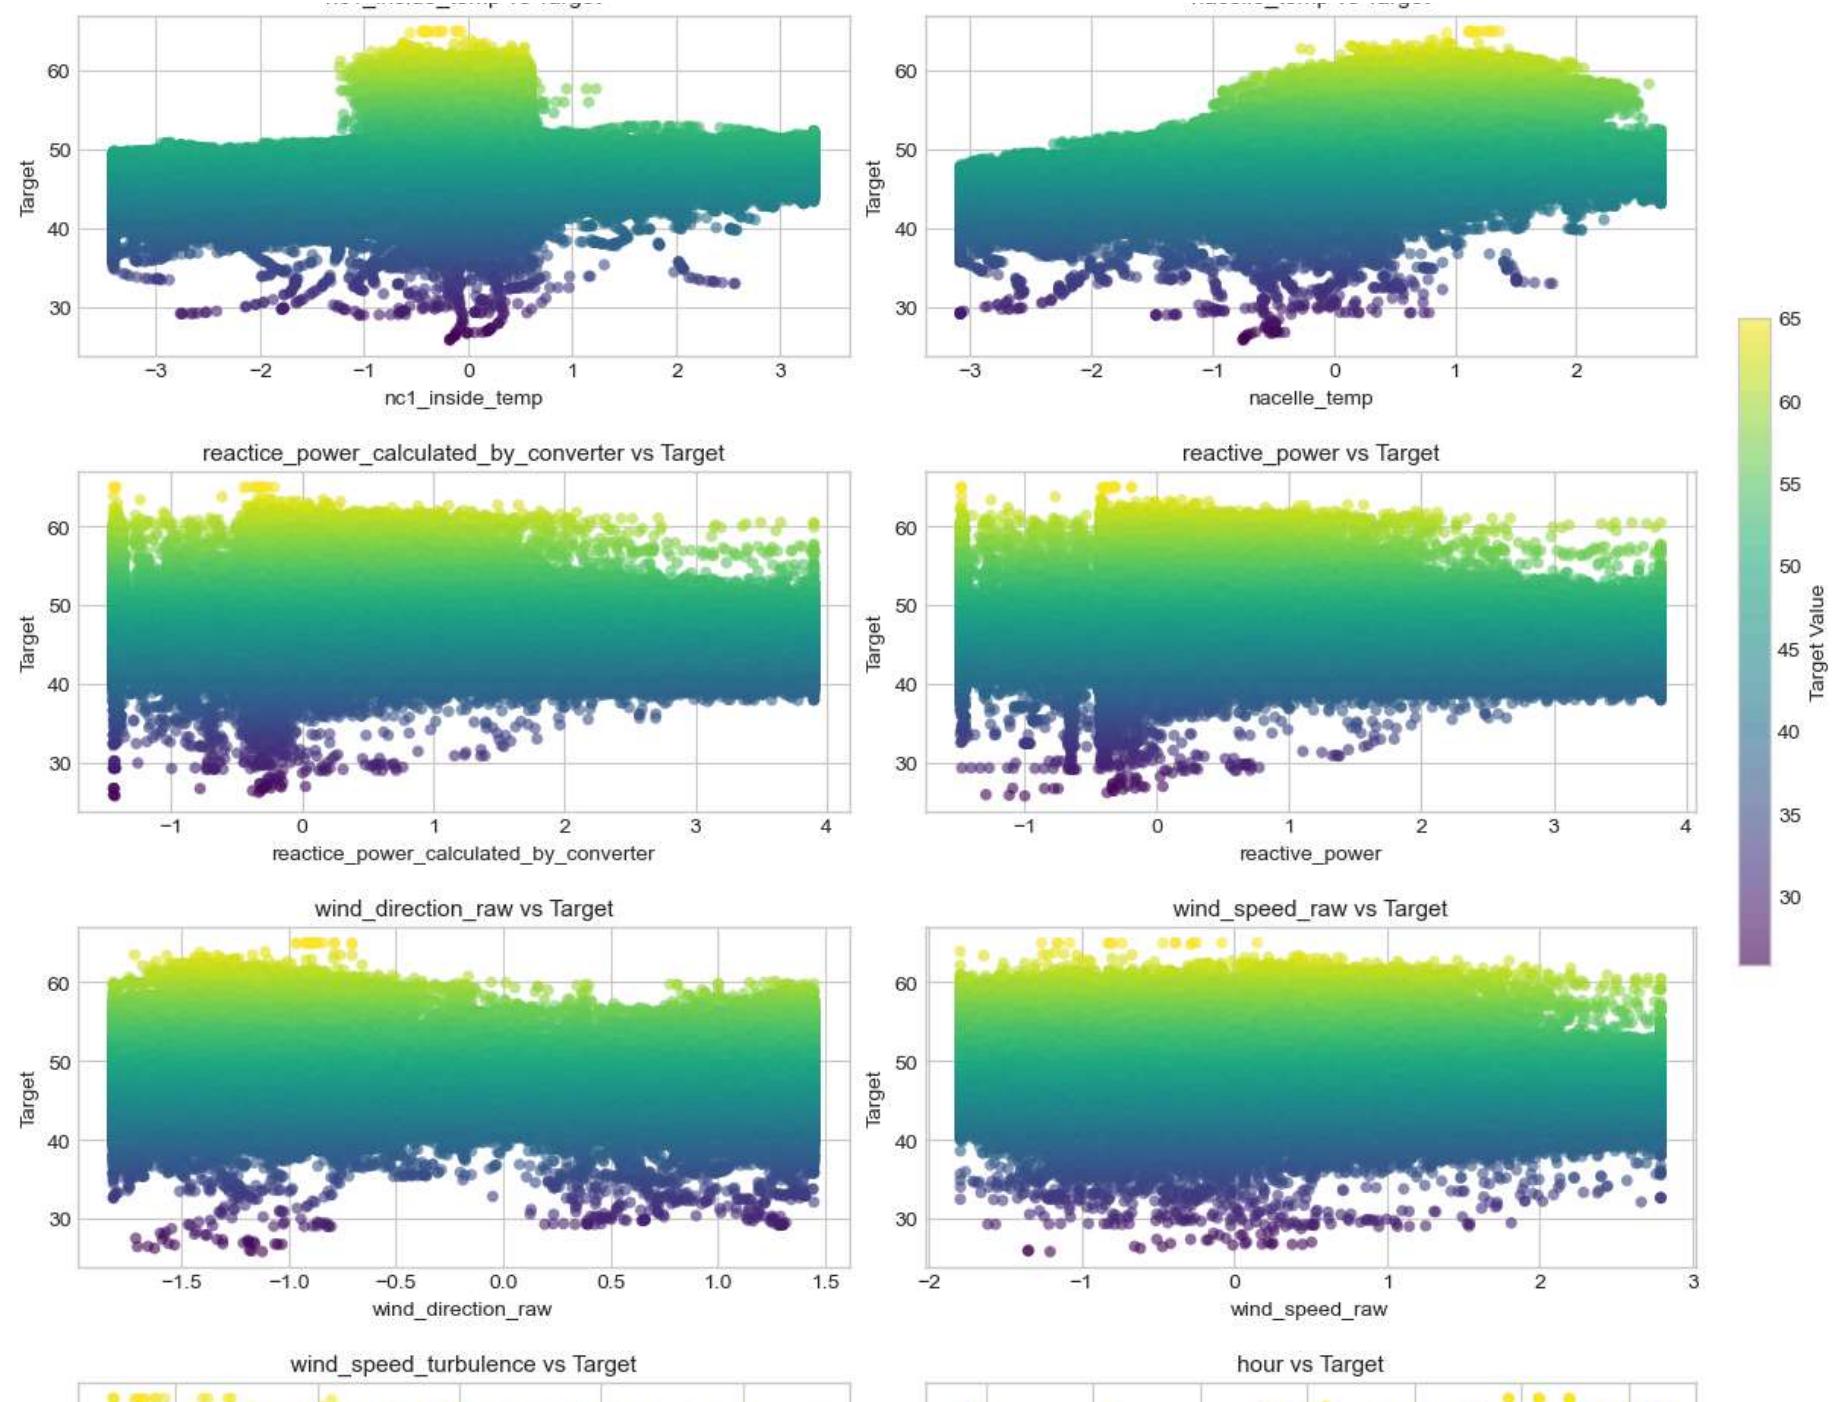
```
        edgecolors='none')
axes[i].set_title(f'{col} vs Target', fontsize=11)
axes[i].set_xlabel(f'{col}\n')
axes[i].set_ylabel('Target')

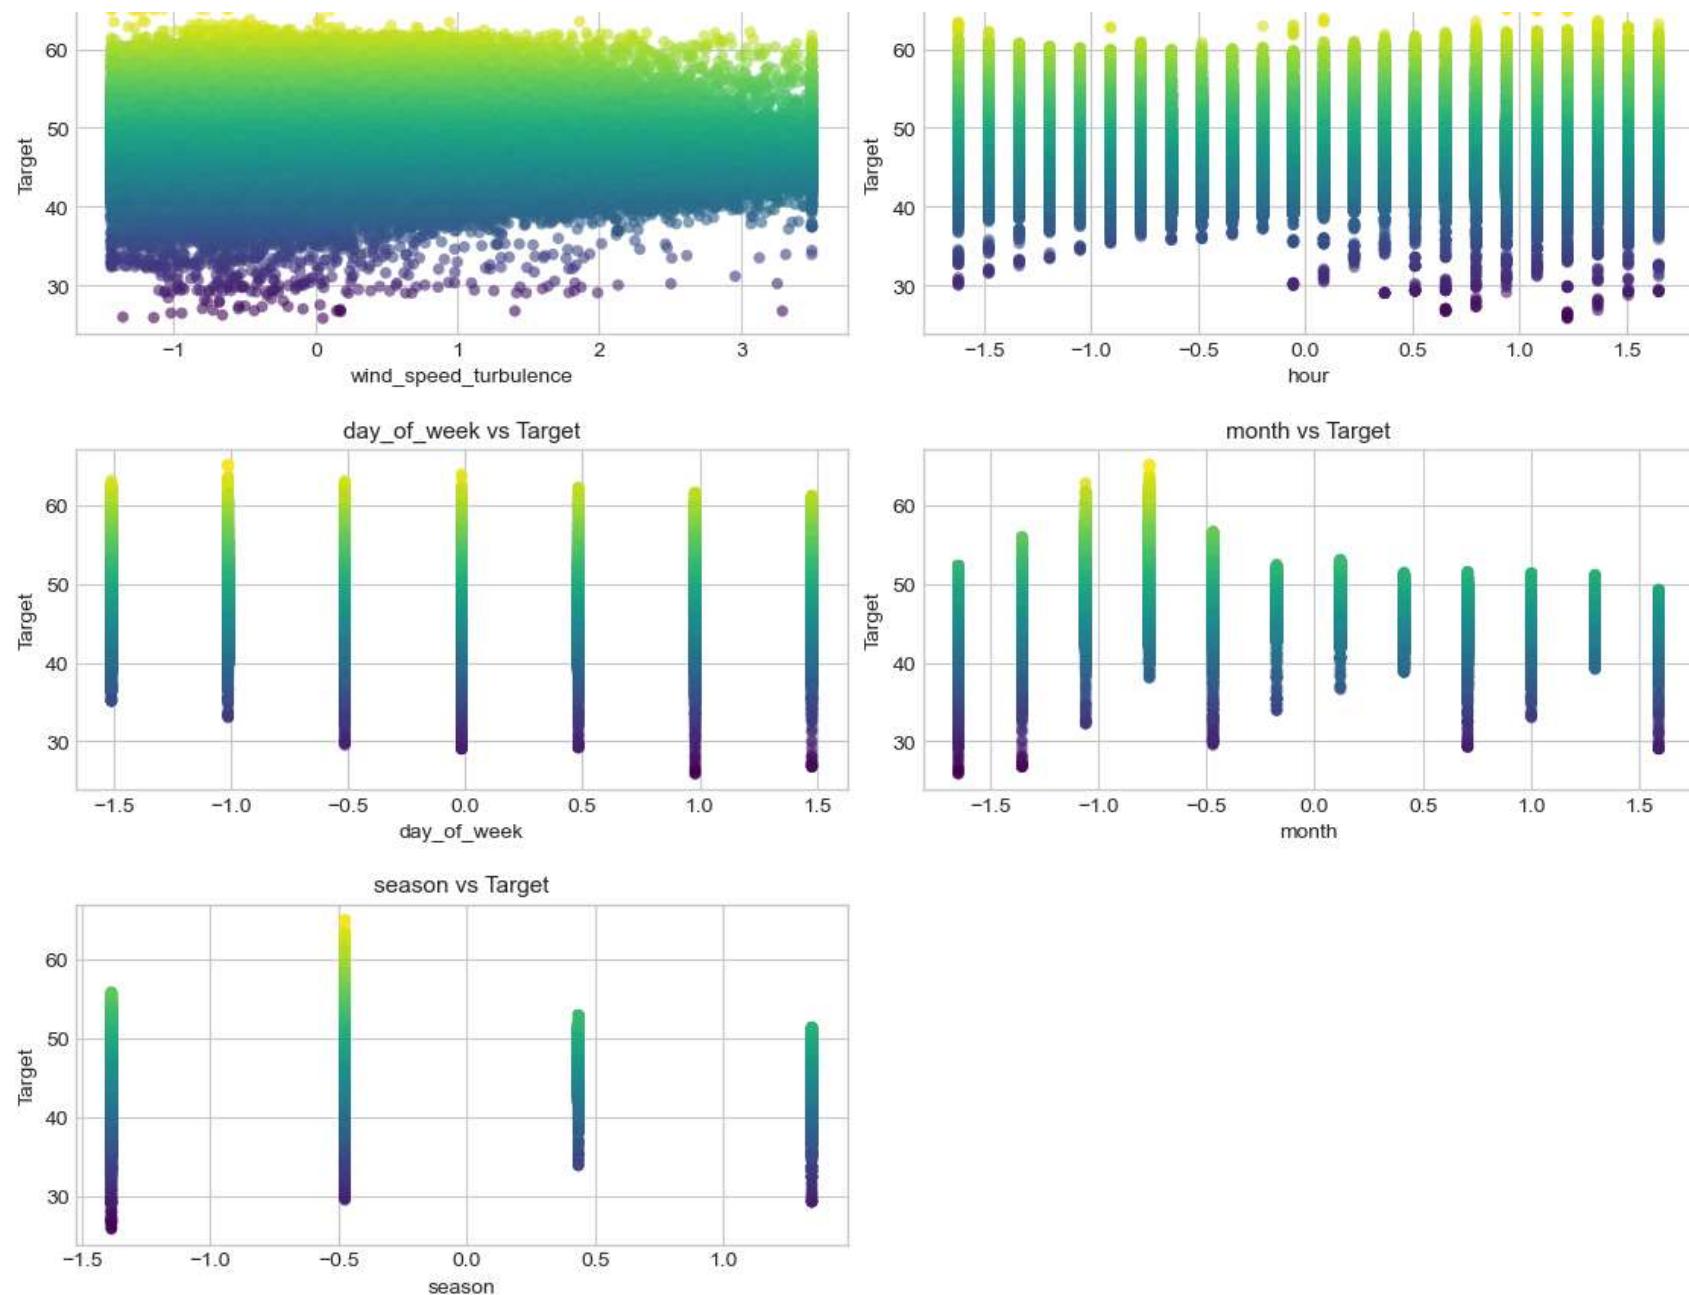
# Add colorbar only once on the right
fig.colorbar(sc, ax=axes, orientation='vertical', fraction=0.02, pad=0.02, label='Target Value')

# Hide any unused subplots
for j in range(i + 1, len(axes)):
    fig.delaxes(axes[j])

# plt.tight_layout()
plt.show()
```

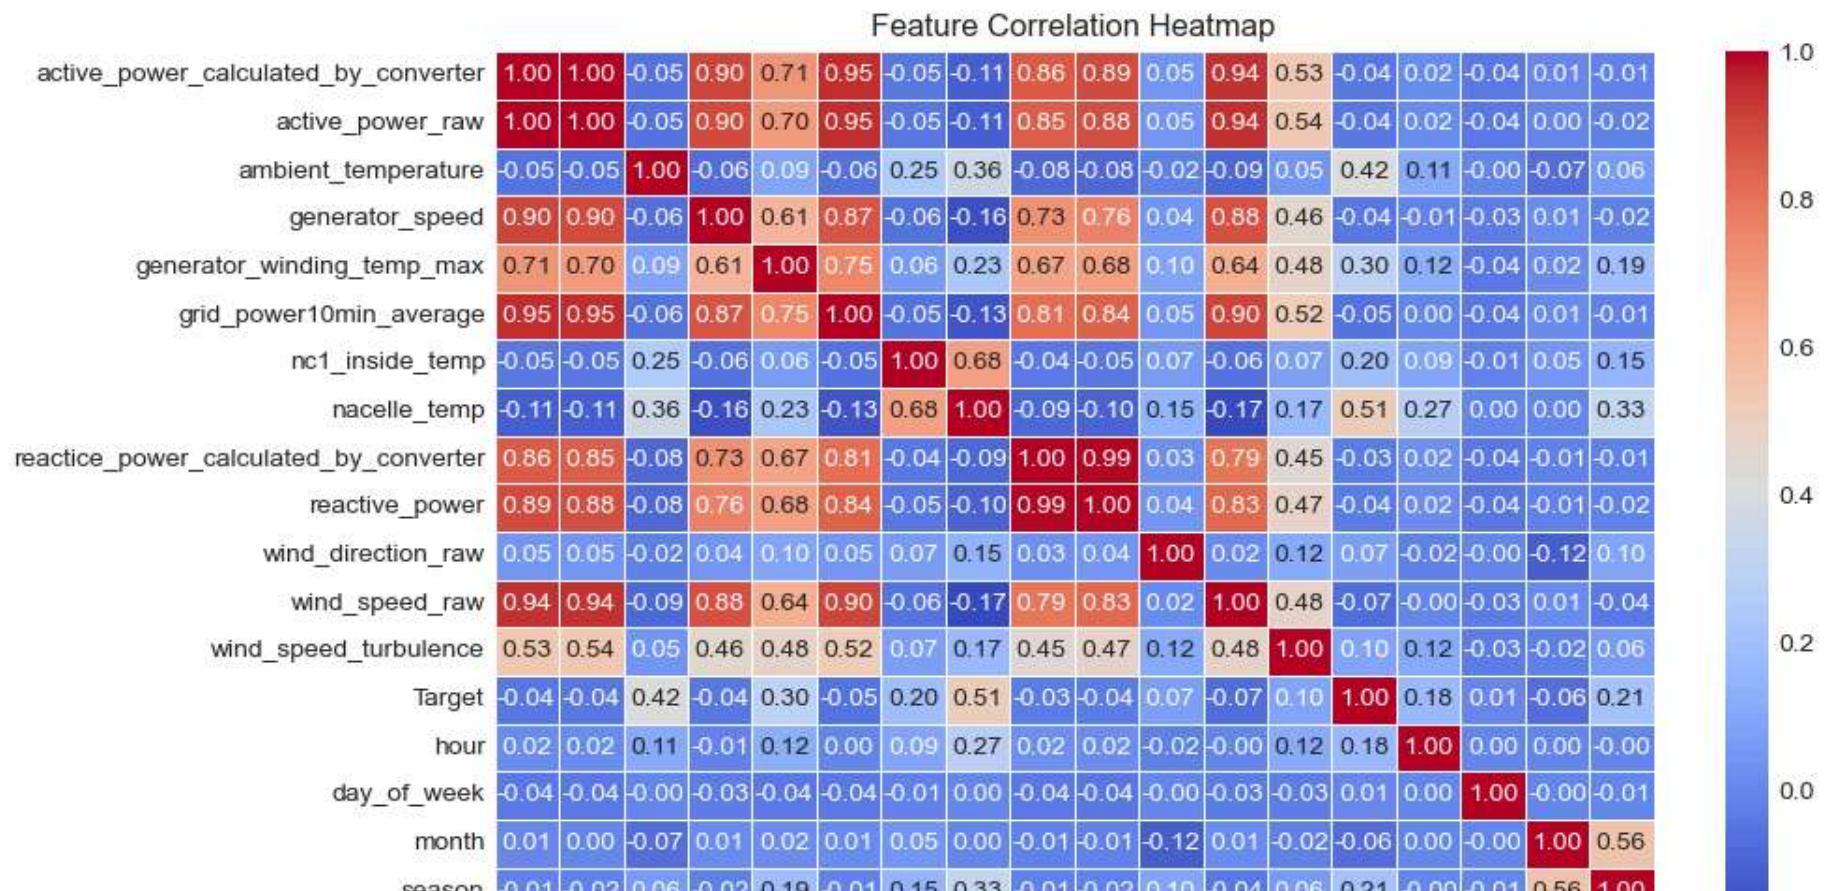






```
In [26]: # Correlation Heatmap
corr_matrix = df.corr()
plt.figure(figsize = (10, 6))
sns.heatmap(corr_matrix, annot = True, cmap = 'coolwarm', fmt = ".2f", linewidths = 0.5)
plt.title("Feature Correlation Heatmap")
plt.show()

target_corr = corr_matrix['Target'].sort_values(ascending = False)
print("\n\nTop features correlated with target:\n",target_corr)
```



active_power_calculated_by_converter	active_power_raw	ambient_temperature	generator_speed	generator_winding_temp_max	grid_power10min_average	nc1_inside_temp	nacelle_temp	reactive_power_calculated_by_converter	reactive_power	wind_direction_raw	wind_speed_raw	wind_speed_turbulence	Target	hour	day_of_week	month	season	
active_power_calculated_by_converter	1.00	1.00	-0.05	0.90	0.71	0.95	-0.05	-0.11	0.86	0.89	0.05	0.94	0.53	-0.04	0.02	-0.04	0.01	-0.01
active_power_raw	1.00	1.00	-0.05	0.90	0.70	0.95	-0.05	-0.11	0.85	0.88	0.05	0.94	0.54	-0.04	0.02	-0.04	0.00	-0.02
ambient_temperature	-0.05	-0.05	1.00	-0.06	0.09	-0.06	0.25	0.36	-0.08	-0.08	-0.02	-0.09	0.05	0.42	0.11	-0.00	-0.07	0.06
generator_speed	0.90	0.90	-0.06	1.00	0.61	0.87	-0.06	-0.16	0.73	0.76	0.04	0.88	0.46	-0.04	-0.01	-0.03	0.01	-0.02
generator_winding_temp_max	0.71	0.70	0.09	0.61	1.00	0.75	0.06	0.23	0.67	0.68	0.10	0.64	0.48	0.30	0.12	-0.04	0.02	0.19
grid_power10min_average	0.95	0.95	-0.06	0.87	0.75	1.00	-0.05	-0.13	0.81	0.84	0.05	0.90	0.52	-0.05	0.00	-0.04	0.01	-0.01
nc1_inside_temp	-0.05	-0.05	0.25	-0.06	0.06	-0.05	1.00	0.68	-0.04	-0.05	0.07	-0.06	0.07	0.20	0.09	-0.01	0.05	0.15
nacelle_temp	-0.11	-0.11	0.36	-0.16	0.23	-0.13	0.68	1.00	-0.09	-0.10	0.15	-0.17	0.17	0.51	0.27	0.00	0.00	0.33
reactive_power_calculated_by_converter	0.86	0.85	-0.08	0.73	0.67	0.81	-0.04	-0.09	1.00	0.99	0.03	0.79	0.45	-0.03	0.02	-0.04	-0.01	-0.01
reactive_power	0.89	0.88	-0.08	0.76	0.68	0.84	-0.05	-0.10	0.99	1.00	0.04	0.83	0.47	-0.04	0.02	-0.04	-0.01	-0.02
wind_direction_raw	0.05	0.05	-0.02	0.04	0.10	0.05	0.07	0.15	0.03	0.04	1.00	0.02	0.12	0.07	-0.02	-0.00	-0.12	0.10
wind_speed_raw	0.94	0.94	-0.09	0.88	0.64	0.90	-0.06	-0.17	0.79	0.83	0.02	1.00	0.48	-0.07	-0.00	-0.03	0.01	-0.04
wind_speed_turbulence	0.53	0.54	0.05	0.46	0.48	0.52	0.07	0.17	0.45	0.47	0.12	0.48	1.00	0.10	0.12	-0.03	-0.02	0.06
Target	-0.04	-0.04	0.42	-0.04	0.30	-0.05	0.20	0.51	-0.03	-0.04	0.07	-0.07	0.10	1.00	0.18	0.01	-0.06	0.21
hour	0.02	0.02	0.11	-0.01	0.12	0.00	0.09	0.27	0.02	0.02	-0.02	-0.00	0.12	0.18	1.00	0.00	0.00	-0.00
day_of_week	-0.04	-0.04	0.00	-0.03	-0.04	-0.04	-0.01	0.00	-0.04	-0.04	-0.00	-0.03	-0.03	0.01	0.00	1.00	-0.00	-0.01
month	0.01	0.00	-0.07	0.01	0.02	0.01	0.05	0.00	-0.01	-0.01	-0.12	0.01	-0.02	-0.06	0.00	-0.00	1.00	0.56
season	-0.01	-0.02	0.06	-0.02	0.19	-0.01	0.15	0.33	-0.01	-0.02	0.10	-0.04	0.06	0.21	-0.00	-0.01	0.56	1.00

Top features correlated with target:

Target	1.000000
nacelle_temp	0.511920
ambient_temperature	0.418106
generator_winding_temp_max	0.297822
season	0.213751
nc1_inside_temp	0.198147
hour	0.176714
wind_speed_turbulence	0.096724
wind_direction_raw	0.070123
day_of_week	0.009359
reactice_power_calculated_by_converter	-0.034431
reactive_power	-0.035796
active_power_raw	-0.040692
active_power_calculated_by_converter	-0.041371
generator_speed	-0.042494
grid_power10min_average	-0.049842
month	-0.063481
wind_speed_raw	-0.069015

Name: Target, dtype: float64

In [27]: *# Checking Multicollinearity with VIF (Variance Inflation Factor)*

```
X = df.copy()

X = X.drop(columns = ['Target'], errors = 'ignore')

vif_data = pd.DataFrame()
vif_data['Feature'] = X.columns
vif_data['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]

vif_data.sort_values(by = "VIF", ascending = False)
```

Out[27]:

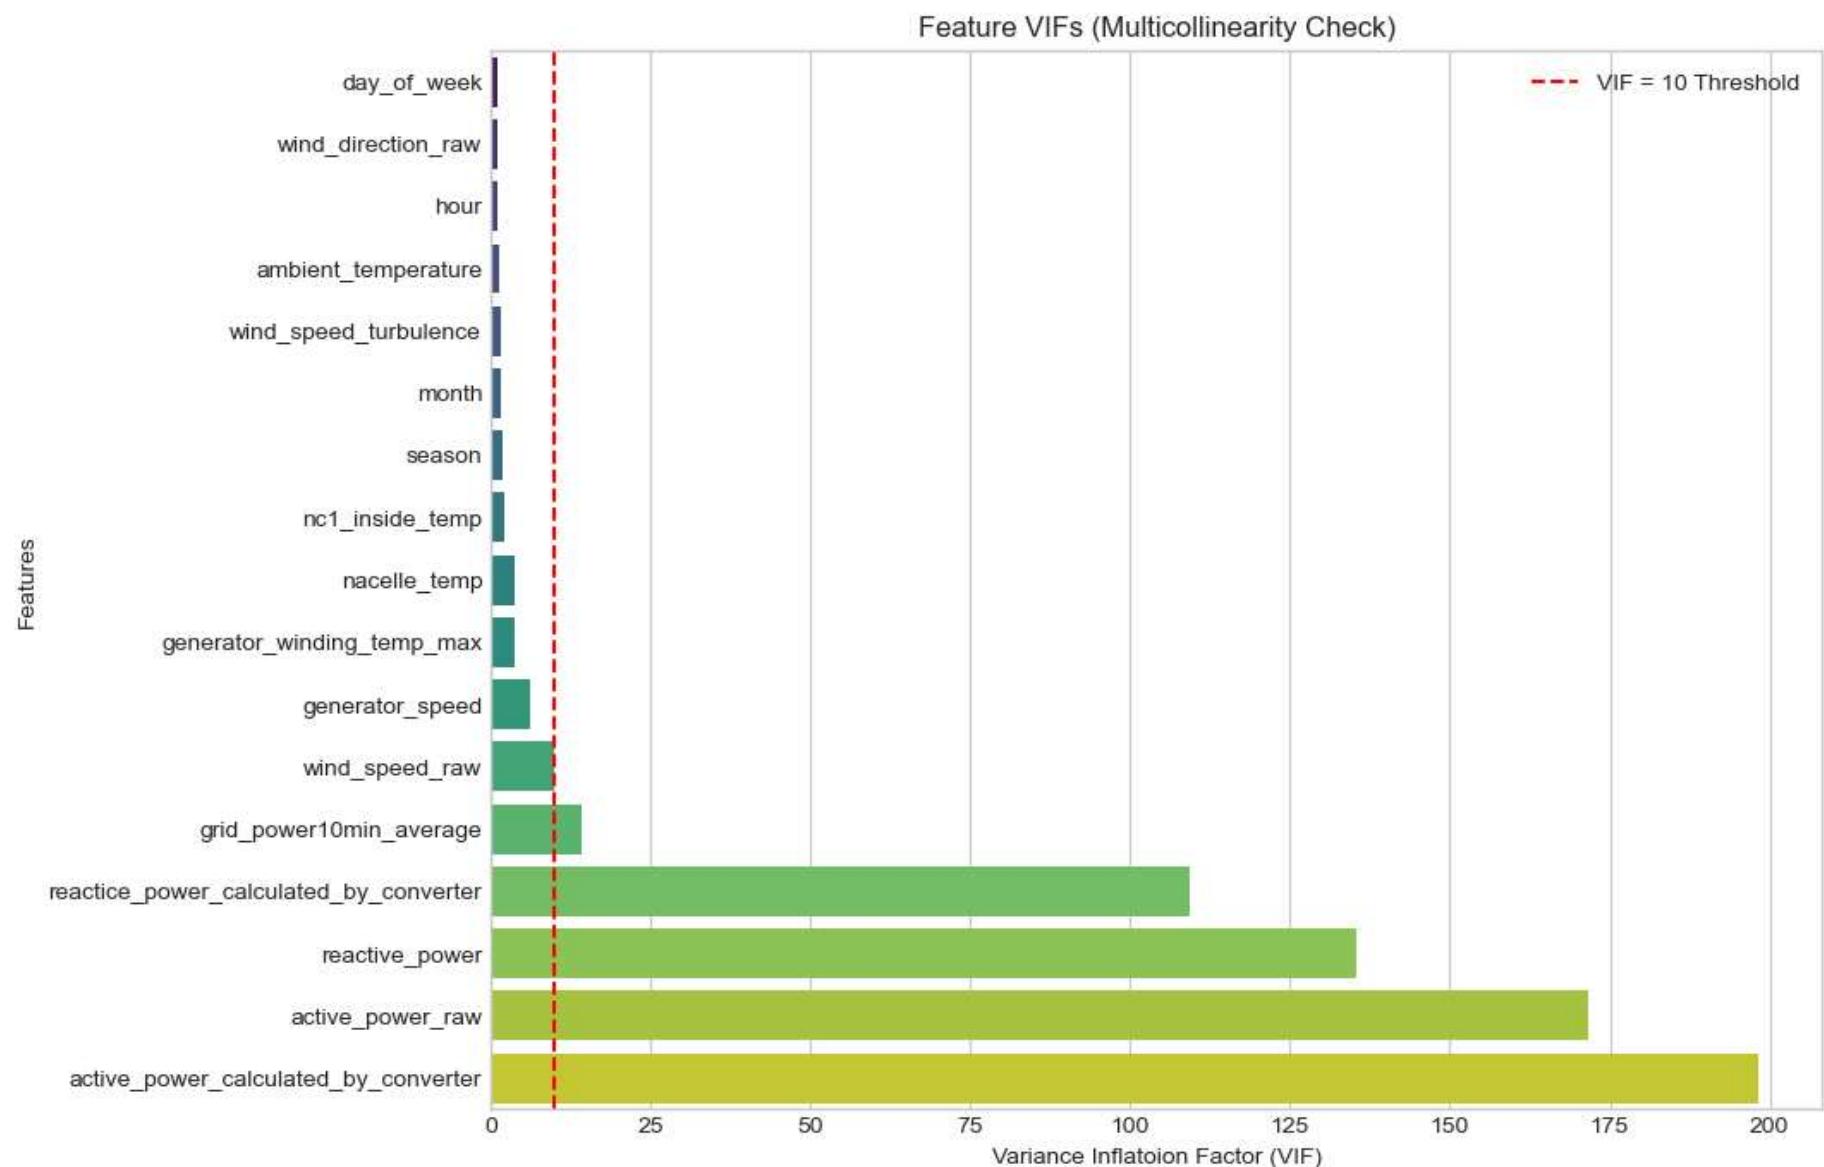
	Feature	VIF
0	active_power_calculated_by_converter	198.305139
1	active_power_raw	171.769409
9	reactive_power	135.326409
8	reactice_power_calculated_by_converter	109.274791
5	grid_power10min_average	14.210131
11	wind_speed_raw	9.872437
3	generator_speed	6.128673
4	generator_winding_temp_max	3.733707
7	nacelle_temp	3.644012
6	nc1_inside_temp	2.203641
16	season	1.888966
15	month	1.645303
12	wind_speed_turbulence	1.563048
2	ambient_temperature	1.182775
13	hour	1.137121
10	wind_direction_raw	1.084663
14	day_of_week	1.002409

In [28]:

```
# VIF Visualization
vif_data_sorted = vif_data.sort_values(by = 'VIF', ascending = True)

plt.figure(figsize = (10, 8))
sns.barplot(x = 'VIF', y = 'Feature', data = vif_data_sorted,hue = 'Feature', palette = 'viridis', legend = False)
plt.axvline(x = 10, color = 'red', linestyle = '--', label = 'VIF = 10 Threshold')
plt.title('Feature VIFs (Multicollinearity Check)')
```

```
plt.xlabel('Variance Inflation Factor (VIF)')
plt.ylabel('Features')
plt.legend()
plt.show()
```



```
In [29]: # Feature Selection
```

```
In [30]: #Features to drop for Linear Regression
X_lr = df.drop(['active_power_raw', 'reactive_power_calculated_by_converter', 'Target'], axis = 1).copy()

# Features to drop for Tree-Based Models
X_tree = df.drop(['Target'], axis = 1).copy()

# Target
y = df['Target']
```

```
In [31]: # One-Hot Encoding the Season feature
if 'season' in X_lr.columns:
    X_lr_ohe = pd.get_dummies(X_lr, columns = ['season'], prefix = 'season', drop_first = True)
else:
    X_lr_ohe = X_lr.copy()
```

```
In [32]: # Features we don't want to scale
time_cols = [c for c in ['hour', 'day_of_week', 'month'] if c in X_lr_ohe.columns]
season_dummy_cols = [c for c in X_lr_ohe.columns if c.startswith('season_')]
dont_scale = time_cols + season_dummy_cols
```

```
In [33]: # Numeric columns to scale
numeric_cols_to_scale = [c for c in X_lr_ohe.columns if c not in dont_scale]
```

```
In [34]: print("Do not Scale (Keep as is):", dont_scale)
print("Numeric columns to scale:", numeric_cols_to_scale)
```

Do not Scale (Keep as is): ['hour', 'day_of_week', 'month', 'season_-0.47581831258795315', 'season_0.4337021794943158', 'season_1.3432226715765847']

Numeric columns to scale: ['active_power_calculated_by_converter', 'ambient_temperature', 'generator_speed', 'generator_winding_temp_max', 'grid_power10min_average', 'nc1_inside_temp', 'nacelle_temp', 'reactive_power', 'wind_direction_raw', 'wind_speed_raw', 'wind_speed_turbulence']

```
In [35]: # Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(X_lr_ohe, y, test_size = 0.20, random_state = 42)
```

```
In [36]: # Scaling Numeric Columns
```

```
scaler = StandardScaler()
```

```
In [37]: # Fit & Transform both train and test- Numeric Features
X_train_scaled = X_train
X_test_scaled = X_test

if numeric_cols_to_scale:
    X_train_scaled[numeric_cols_to_scale] = scaler.fit_transform(X_train[numeric_cols_to_scale])
    X_test_scaled[numeric_cols_to_scale] =scaler.fit_transform(X_test[numeric_cols_to_scale])
```

```
In [38]: print("\nShapes:")
print("X_train:", X_train_scaled.shape, "\nX_test:", X_test_scaled.shape)
print("\ny_train:", y_train.shape, "\ny_test:", y_test.shape )
```

Shapes:
X_train: (727683, 17)
X_test: (181921, 17)

y_train: (727683,)
y_test: (181921,)

```
In [39]: print("\nSample (First 3 Rows) of scaled training data:")
display(X_train_scaled.head(3))
```

Sample (First 3 Rows) of scaled training data:

	active_power_calculated_by_converter	ambient_temperature	generator_speed	generator_winding_temp_max	grid_power10min_average
764808	1.176605	0.877374	1.437085	2.781710	2.52193
565905	-0.752966	0.616338	-0.903765	-0.324972	-0.46018
459081	-0.422716	0.557758	-0.441565	-0.291983	-0.28586



```
In [40]: # --- REDUNDANCY & UNIQUENESS CHECK ---
```

```
# 1) Check exact duplicates (across all columns)
total_rows = len(df)
duplicate_rows = df.duplicated().sum()
```

```
print(f"Total Rows: {total_rows:,}")
print(f"Exact Duplicate Rows: {duplicate_rows:,}")
print(f"Duplicate Percentage: {(duplicate_rows / total_rows) * 100:.2f}%")

# 2) Drop duplicates to get unique data for further comparison
df_unique = df.drop_duplicates()
unique_rows = len(df_unique)

print(f"\nUnique Rows: {unique_rows:,}")
print(f"Unique Data Percentage: {(unique_rows / total_rows) * 100:.2f}%")

# 3) Check near-duplicates - same readings on key columns (non-random)
# We'll use a few key physical features
key_features = ['ambient_temperature', 'nacelle_temp', 'generator_speed', 'wind_speed', 'Target']
key_features = [f for f in key_features if f in df.columns]

if key_features:
    df_key = df[key_features]
    near_duplicates = df_key.duplicated().sum()
    print(f"\nNear-Duplicate Rows (based on key turbine readings): {near_duplicates:,}")
    print(f"Near-Duplicate Percentage: {(near_duplicates / total_rows) * 100:.2f}%")
else:
    print("\nNo matching key features found for near-duplicate check.")

# 4) Optional: simple compression ratio (for intuition)
compression_ratio = unique_rows / total_rows
print(f"\nData Redundancy Ratio (1 = no redundancy, 0.5 = 50% repetitive): {compression_ratio:.3f}")
```

Total Rows: 909,604
Exact Duplicate Rows: 0
Duplicate Percentage: 0.00%

Unique Rows: 909,604
Unique Data Percentage: 100.00%

Near-Duplicate Rows (based on key turbine readings): 0
Near-Duplicate Percentage: 0.00%

Data Redundancy Ratio (1 = no redundancy, 0.5 = 50% repetitive): 1.000

```
In [41]: # *****Linear Regression Baseline Model*****
```

```
In [42]: # Linear Regression Baseline Model
lr_model = LinearRegression()
```

```
In [43]: lr_model.fit(X_train_scaled, y_train)
```

```
Out[43]: ▾ LinearRegression ⓘ ?
```

```
LinearRegression()
```

```
In [44]: # Predict on test data
y_pred_lr = lr_model.predict(X_test_scaled)
```

```
In [45]: # Evaluate Performance
mae_lr = mean_absolute_error(y_test, y_pred_lr)
rmse_lr = np.sqrt(mean_squared_error(y_test, y_pred_lr))
r2_lr = r2_score(y_test, y_pred_lr)
```

```
In [46]: # Adjusted R2
n = len(y_test)
p = X_test_scaled.shape[1]
adj_r2_lr = 1 - (1 - r2_lr) * (n - 1) / (n - p - 1)
```

```
In [47]: print(f"Linear Regression Performance on Full Data:")
print("-----")
print(f"MAE : {mae_lr:.4f}")
print(f"RMSE : {rmse_lr:.4f}")
print(f"R2 : {r2_lr:.4f}")
print(f"Adjusted R2 : {adj_r2_lr:.4f}")
```

Linear Regression Performance on Full Data:

MAE : 1.3793

RMSE : 1.9361

R² : 0.4511

Adjusted R² : 0.4510

```
In [48]: # Define tolerance (within ±10% of actual value)
tolerance = 0.10
within_tolerance = np.abs((y_test - y_pred_lr) / y_test) <= tolerance
accuracy = within_tolerance.mean() * 100

print(f"Approximate Accuracy (within ±10% tolerance): {accuracy:.2f}%")
```

Approximate Accuracy (within ±10% tolerance): 97.48%

```
In [49]: # Influential features (Coefficients)
coef_df = pd.DataFrame({
    'Feature': X_train_scaled.columns,
    'Coefficient': lr_model.coef_
}).sort_values(by='Coefficient', key=abs, ascending=False)

print("\nTop 10 Most Influential Features:")
display(coef_df.head(10))
```

Top 10 Most Influential Features:

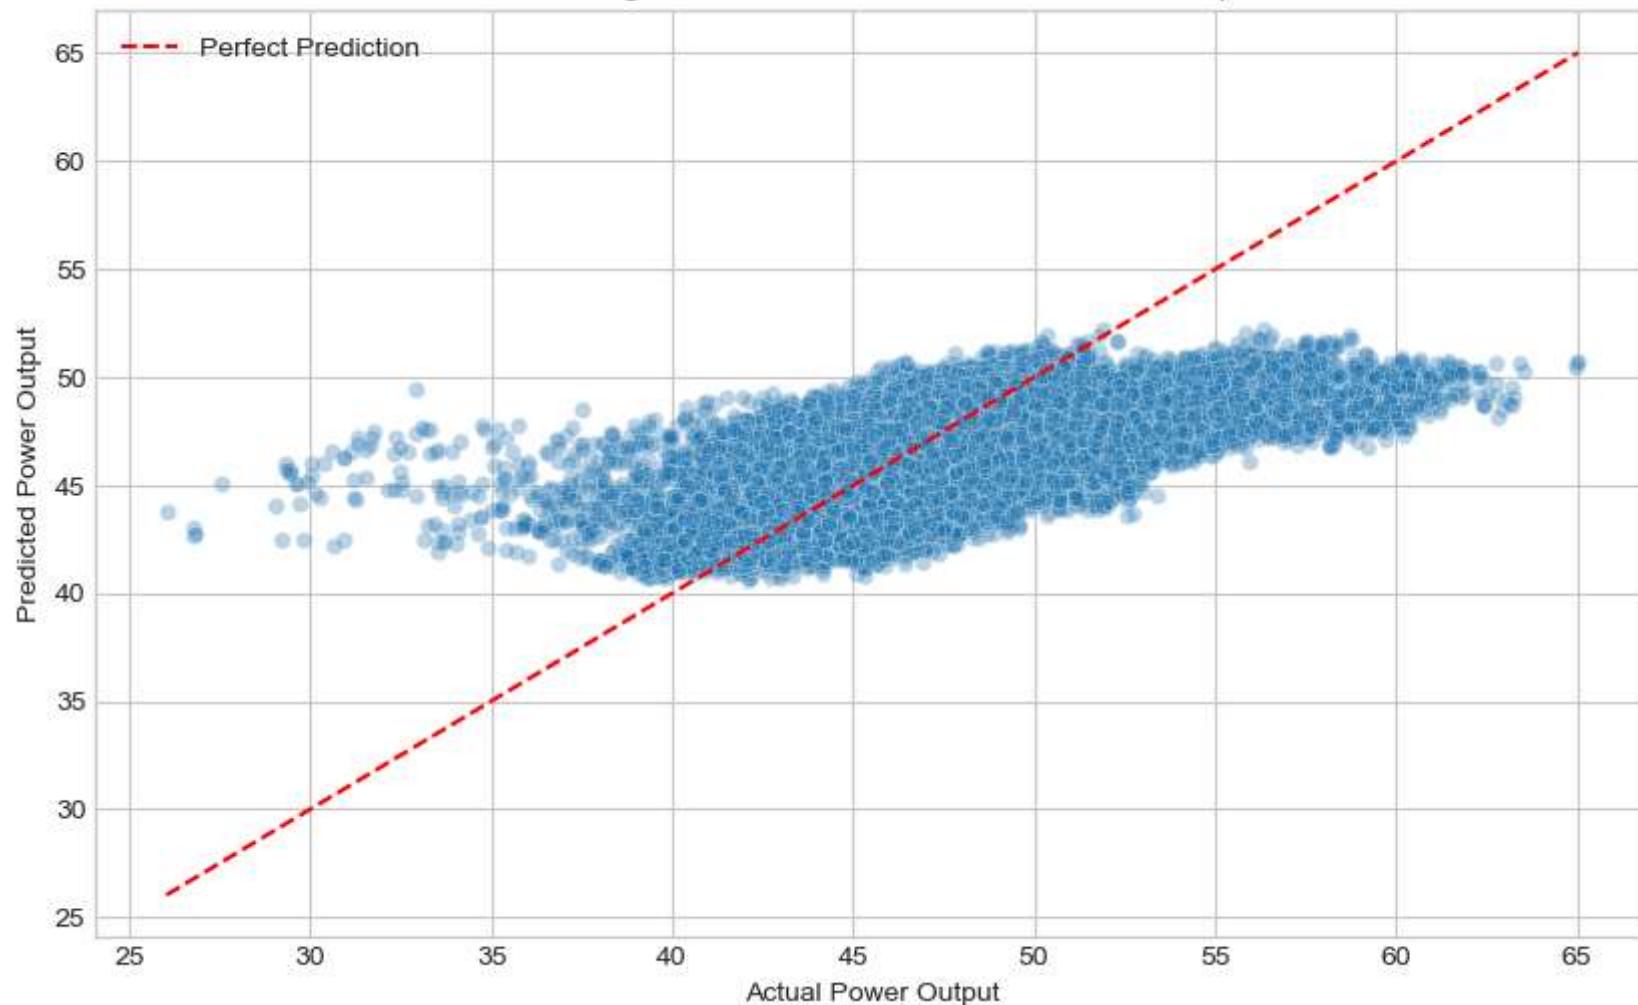
	Feature	Coefficient
14	season_-0.47581831258795315	1.449415
3	generator_winding_temp_max	0.991736
15	season_0.4337021794943158	0.921972
4	grid_power10min_average	-0.911043
16	season_1.3432226715765847	0.780385
6	nacelle_temp	0.747685
1	ambient_temperature	0.666495
5	nc1_inside_temp	-0.344781
2	generator_speed	0.306365
0	active_power_calculated_by_converter	-0.187417

```
In [50]: # Convert predictions to Series aligned with y_test
y_pred_series = pd.Series(y_pred_lr, index=y_test.index)

# Actual vs Predicted
plt.figure(figsize=(10,6))
sns.scatterplot(x=y_test, y=y_pred_series, alpha=0.3) # alpha=0.3 to handle dense points
plt.plot([y_test.min(), y_test.max()],
          [y_test.min(), y_test.max()],
          color='red', linestyle='--', label='Perfect Prediction')
plt.xlabel("Actual Power Output")
plt.ylabel("Predicted Power Output")
plt.title("Linear Regression: Actual vs Predicted Power Output")
plt.legend()
plt.savefig("scatter_actual_vs_prediction.png", dpi=300, bbox_inches="tight")

plt.show()
```

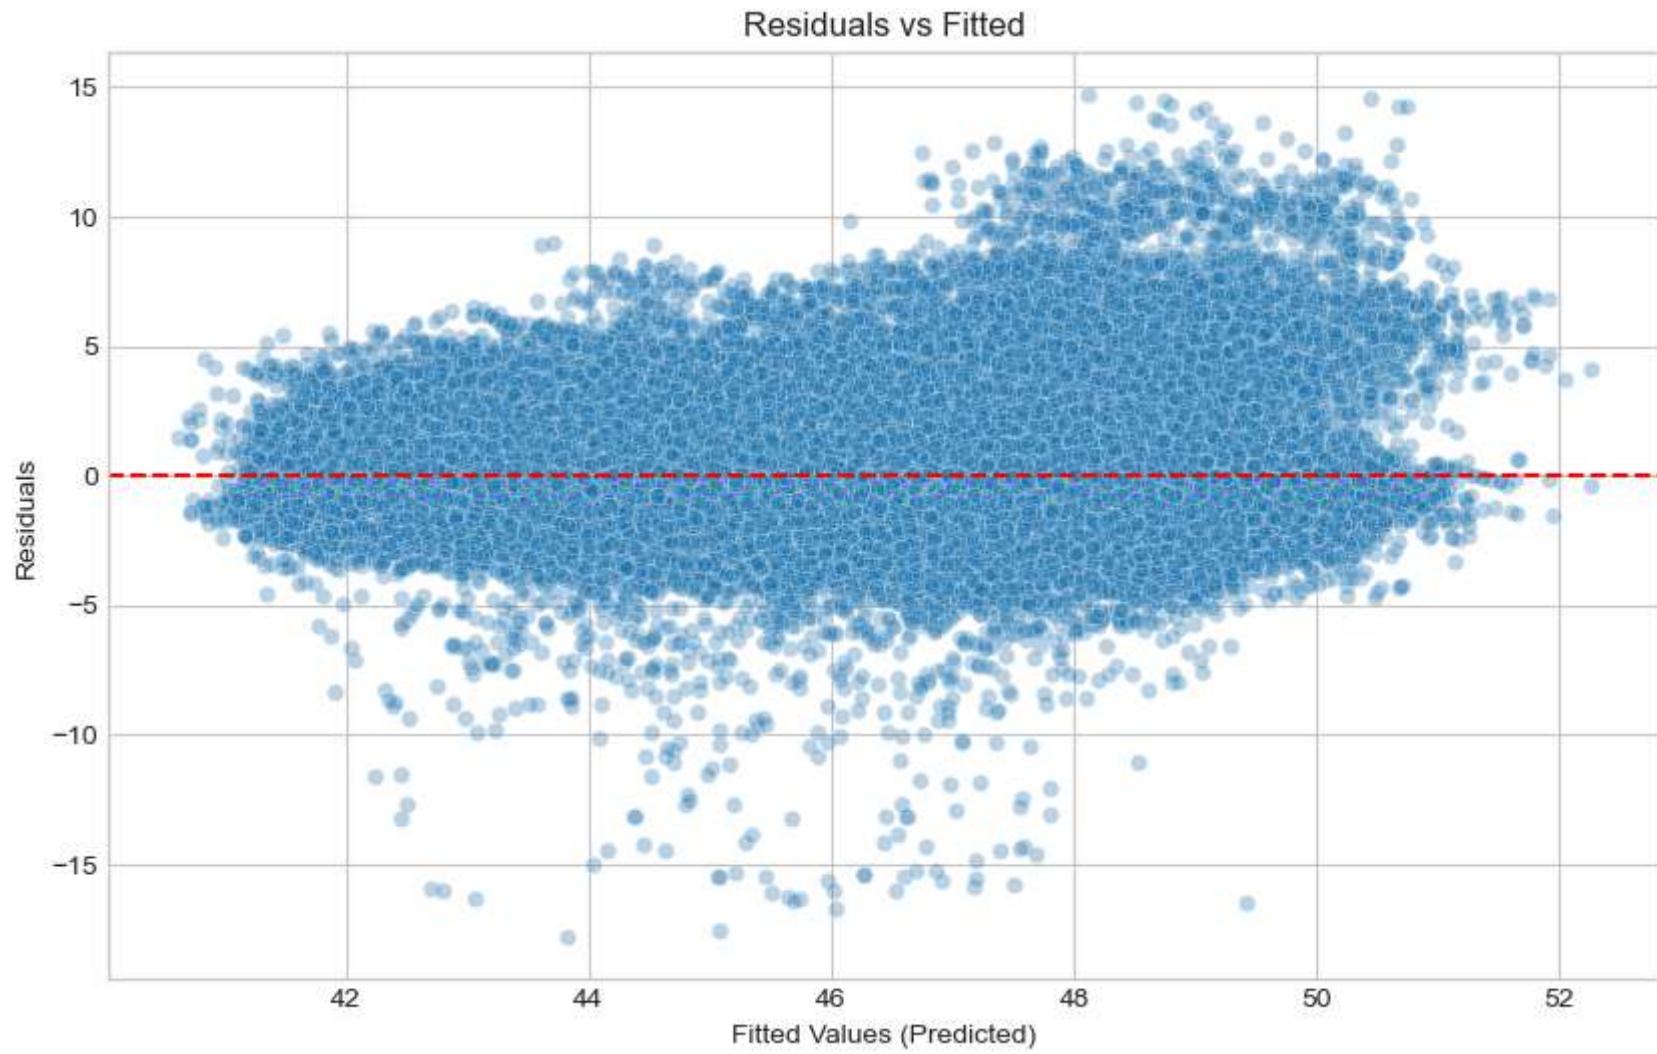
Linear Regression: Actual vs Predicted Power Output



```
In [51]: # Convert predictions to Series aligned with y_test
y_pred_series = pd.Series(y_pred_lr, index=y_test.index)
residuals = y_test - y_pred_series

# Residuals vs Fitted
plt.figure(figsize=(10,6))
sns.scatterplot(x=y_pred_series, y=residuals, alpha=0.3)
plt.axhline(0, color='red', linestyle='--')
```

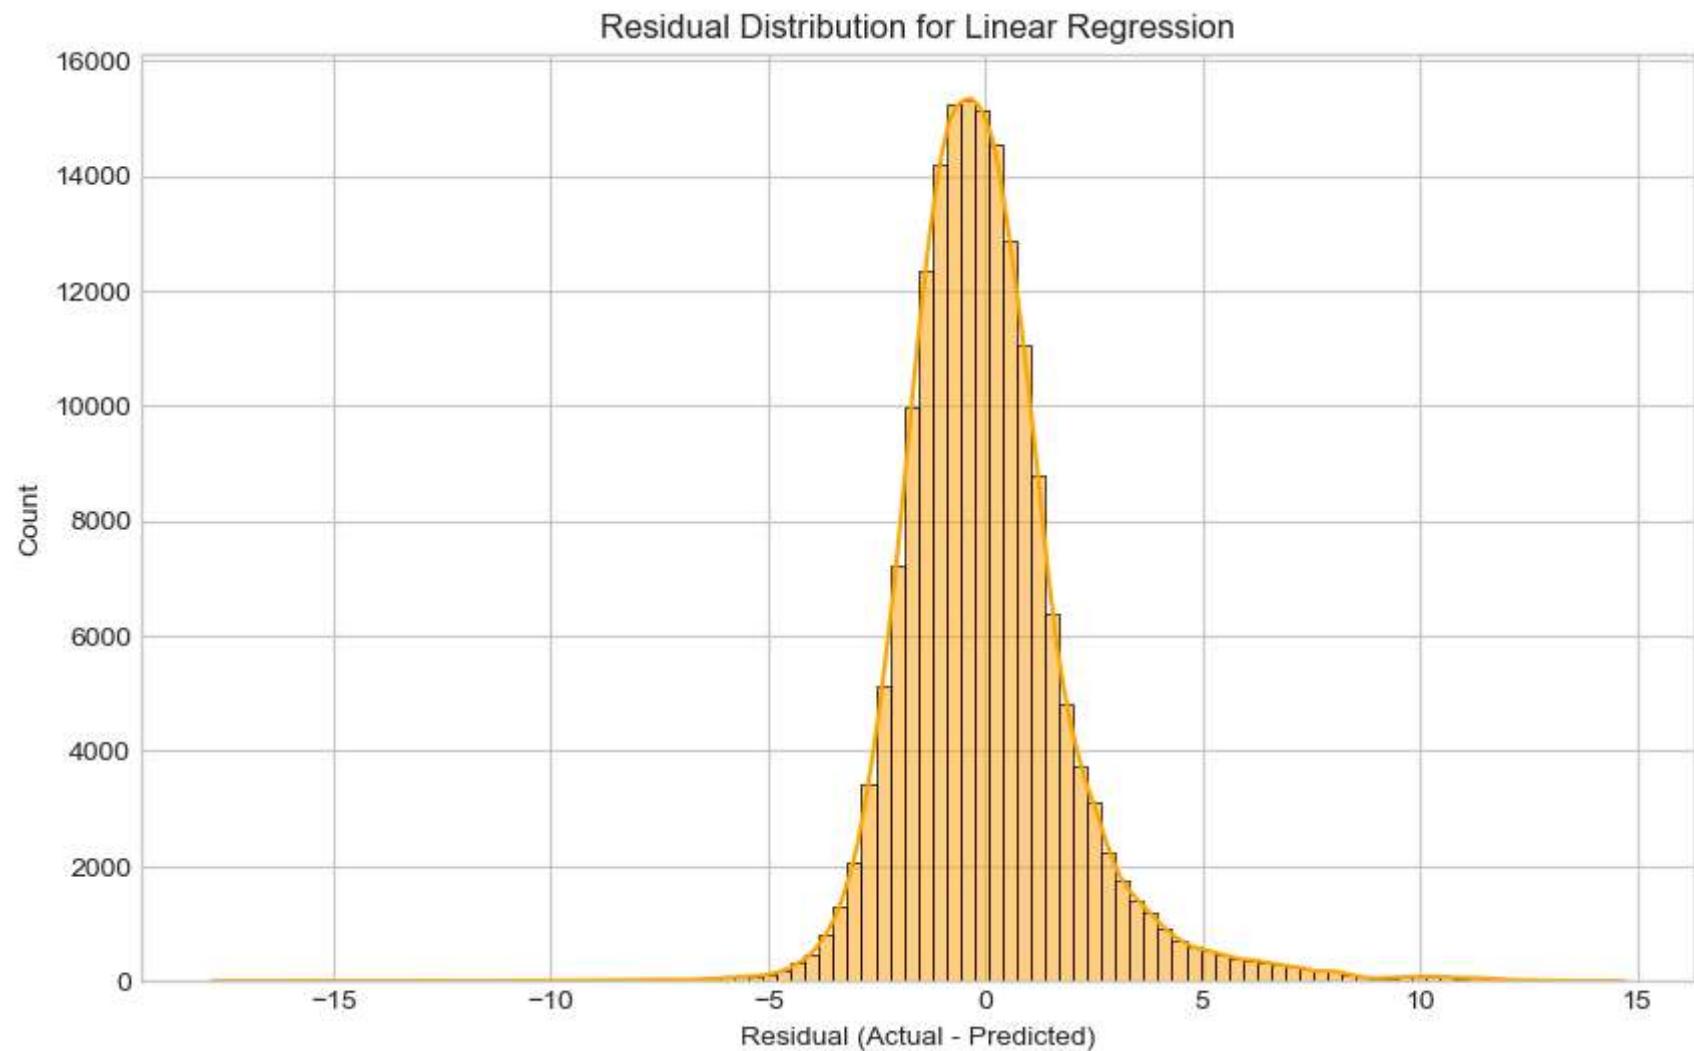
```
plt.xlabel("Fitted Values (Predicted)")  
plt.ylabel("Residuals")  
plt.title("Residuals vs Fitted")  
plt.show()
```



```
In [52]: # Residuals plot: Prediction errors  
residuals = y_test - y_pred_series  
plt.figure(figsize=(10,6))
```

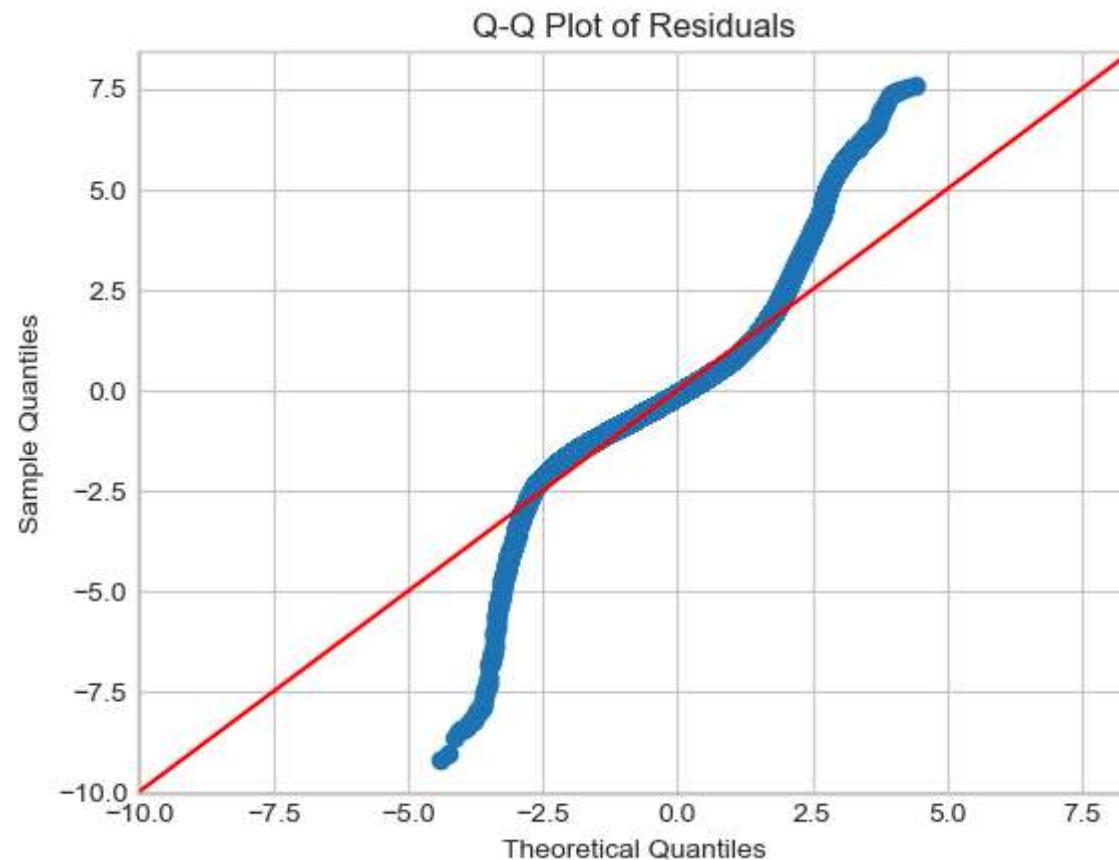
```
sns.histplot(residuals, bins=100, kde=True, color='orange')
plt.xlabel("Residual (Actual - Predicted)")
plt.title("Residual Distribution for Linear Regression")
plt.savefig("Residual distribution.png", dpi=300, bbox_inches="tight")

plt.show()
```



```
In [53]: # Q-Q Plot (Normality check)

sm.qqplot(residuals, line='45', fit=True)
plt.title("Q-Q Plot of Residuals")
plt.show()
```



```
In [54]: X_clean = X_lr_ohe.copy()

# Convert booleans to integers (True->1, False->0)
bool_cols = [c for c in X_clean.columns if X_clean[c].dtype == 'bool']
if bool_cols:
    X_clean[bool_cols] = X_clean[bool_cols].astype(int)
```

```
# Force all columns to numeric (safety) – non-convertible -> NaN
X_clean = X_clean.apply(pd.to_numeric, errors='coerce')
```

In [55]:

```
# Align y and drop rows with any NaN in X or y
y_clean = pd.to_numeric(y, errors='coerce')
mask = X_clean.notna().all(axis=1) & y_clean.notna()
X_clean = X_clean.loc[mask].astype(float)
y_clean = y_clean.loc[mask].astype(float)
```

In [56]:

```
print("After cleaning: X shape =", X_clean.shape, " y shape =", y_clean.shape)
print("Example dtypes:\n", X_clean.dtypes.head(10))
```

After cleaning: X shape = (909604, 17) y shape = (909604,)

Example dtypes:

active_power_calculated_by_converter	float64
ambient_temperature	float64
generator_speed	float64
generator_winding_temp_max	float64
grid_power10min_average	float64
nc1_inside_temp	float64
nacelle_temp	float64
reactive_power	float64
wind_direction_raw	float64
wind_speed_raw	float64
dtype:	object

In [57]:

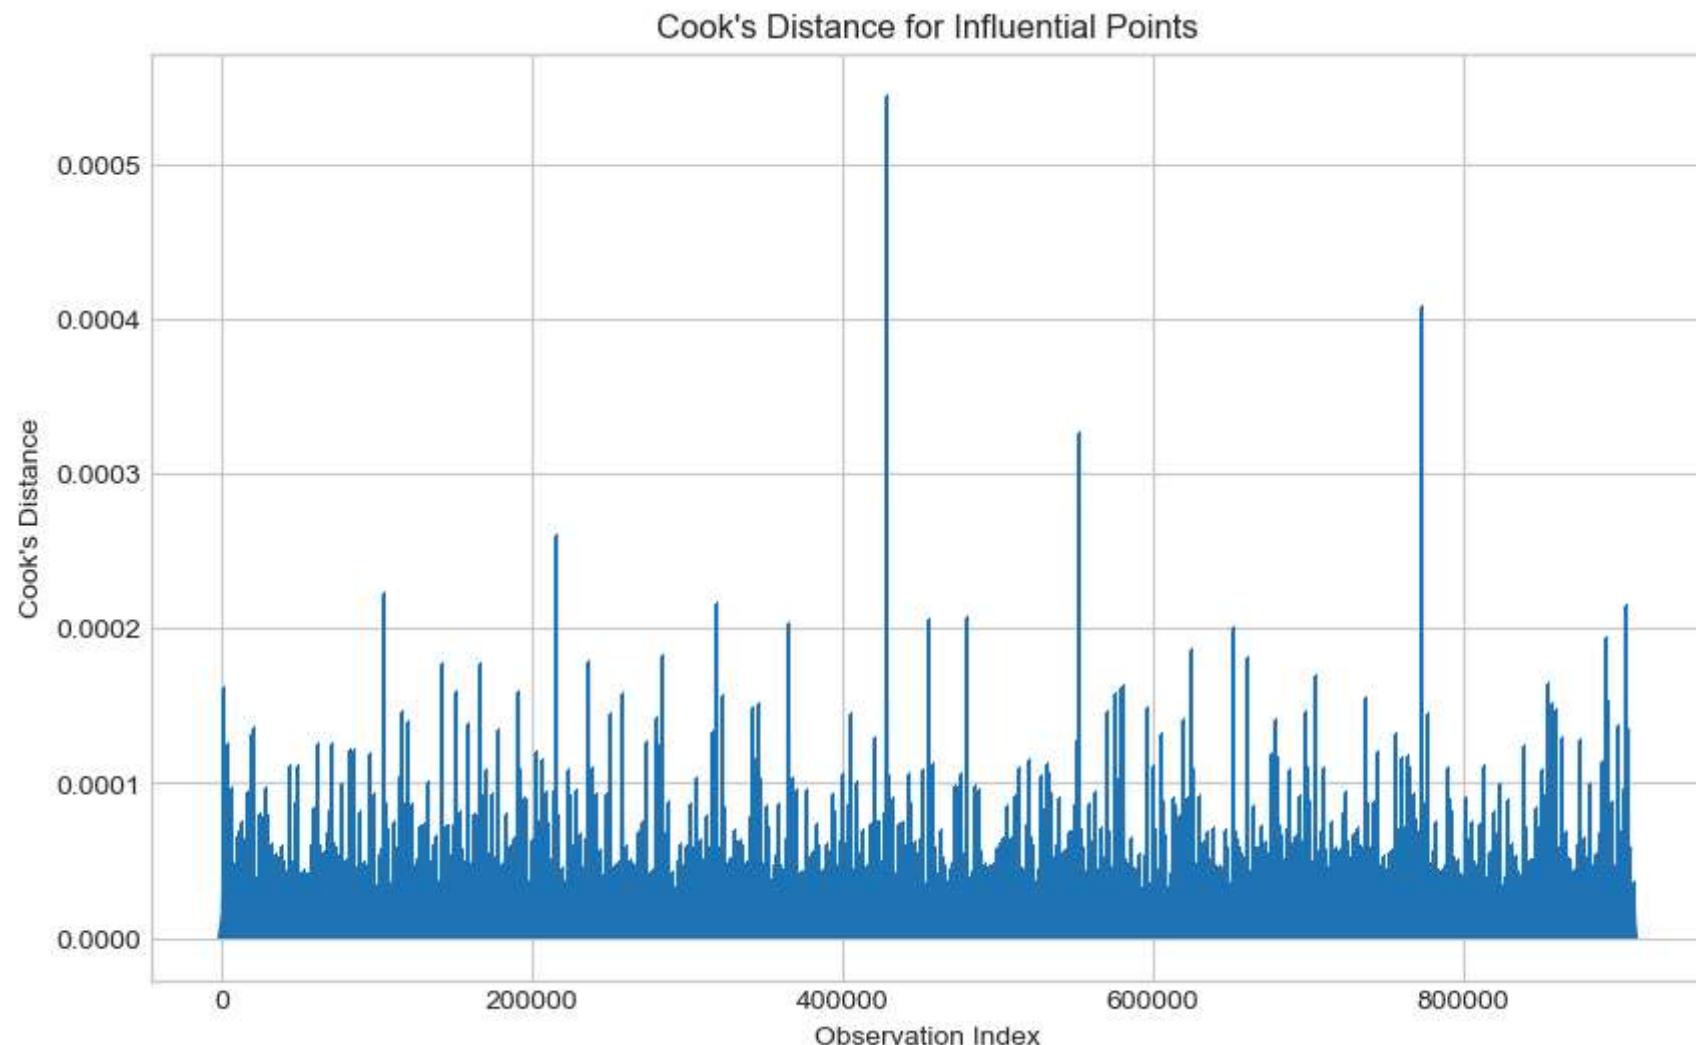
```
X_const = sm.add_constant(X_clean, has_constant='add')
model_sm = sm.OLS(y_clean, X_const).fit()
```

In [58]:

```
# Cook's distance (diagnostics)
influence = model_sm.get_influence()
c, p = influence.cooks_distance

plt.figure(figsize=(10,6))
plt.stem(np.arange(len(c)), c, markerfmt=",", basefmt=" ")
plt.xlabel("Observation Index")
plt.ylabel("Cook's Distance")
plt.title("Cook's Distance for Influential Points")
plt.show()
```

```
# Identify high-influence points using common threshold 4/n
n = len(y_clean)
infl_pts = np.where(c > 4/n)[0]
print("Number of highly influential points:", len(infl_pts))
```



Number of highly influential points: 45674

In [59]: # *****RANDOM FOREST REGRESSOR*****

```
In [60]: # Random Forest Regressor Model
```

```
rf_model = RandomForestRegressor(  
    n_estimators = 100,  
    random_state = 42,  
    n_jobs = -1  
)
```

```
In [61]: # Train the model
```

```
rf_model.fit(X_train, y_train)
```

```
Out[61]:
```

```
RandomForestRegressor  
RandomForestRegressor(n_jobs=-1, random_state=42)
```

```
In [62]: # Predict on test data
```

```
y_pred_rf = rf_model.predict(X_test)
```

```
In [63]: # Evaluate the performance of the Random Forest Model
```

```
mae_rf = mean_absolute_error(y_test, y_pred_rf)  
rmse_rf = np.sqrt(mean_squared_error(y_test, y_pred_rf))  
r2_rf = r2_score(y_test, y_pred_rf)
```

```
In [64]: print("****Random Forest Regressor Results****")  
print(f"Mean Absolute Error (MAE): {mae_rf:.4f}")  
print(f"Root Mean Squared Error (RMSE): {rmse_rf:.4f}")  
print(f"R² Score: {r2_rf:.4f}")
```

****Random Forest Regressor Results****

Mean Absolute Error (MAE): 0.4329

Root Mean Squared Error (RMSE): 0.7609

R² Score: 0.9152

```
In [65]: comparison = pd.DataFrame({
```

```
    "Model": ["Linear Regression", "Random Forest"],  
    "MAE": [mae_lr, mae_rf],  
    "RMSE": [rmse_lr, rmse_rf],
```

```

    "R2": [r2_lr, r2_rf]
})

print("\nModel Comparison:\n")
display(comparison)

```

Model Comparison:

	Model	MAE	RMSE	R ²
0	Linear Regression	1.379340	1.936138	0.451066
1	Random Forest	0.432881	0.760886	0.915221

```
In [66]: # Feature importance scores
importances = rf_model.feature_importances_
feature_names = X_train.columns
```

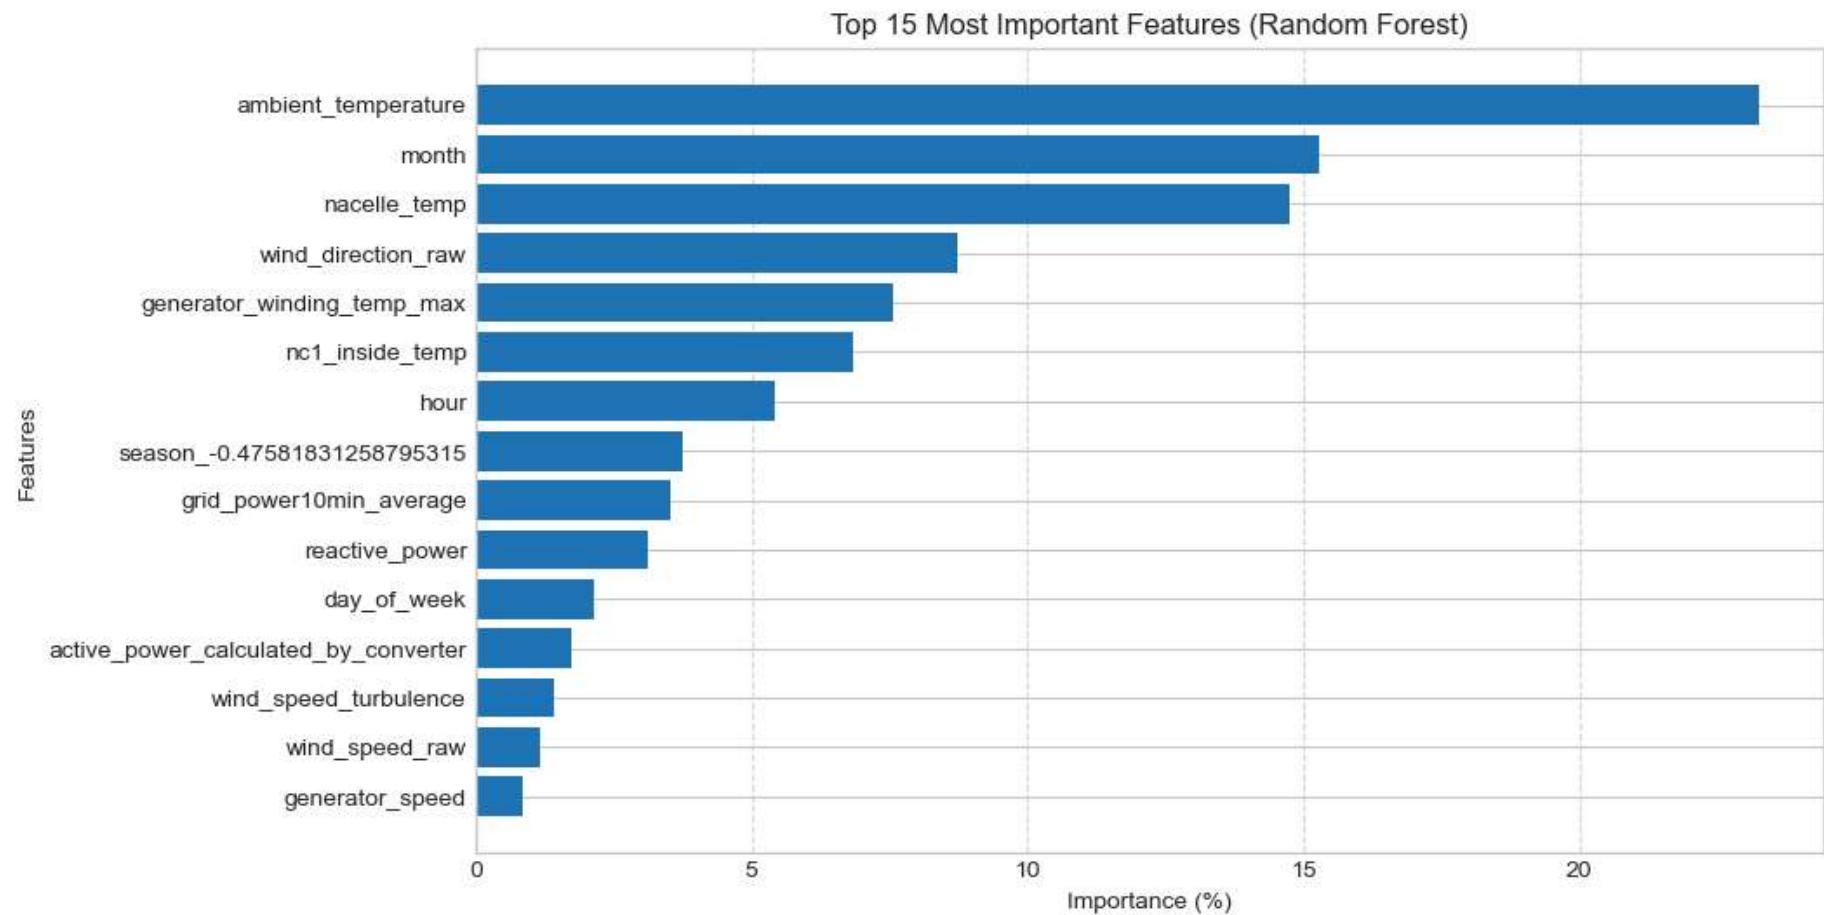
```
In [67]: feat_importance_df = pd.DataFrame({
    'Feature': feature_names,
    'Importance': importances
}).sort_values(by='Importance', ascending=False)
```

```
In [68]: # Normalize importance for percentage scale
feat_importance_df['Importance (%)'] = 100 * feat_importance_df['Importance'] / feat_importance_df['Importance'].sum()
```

```
In [69]: # Display top 15 features
top_features = feat_importance_df.head(15)
```

```
In [70]: # --- Visualization ---
plt.figure(figsize=(10, 6))
plt.barh(top_features['Feature'], top_features['Importance (%)'])
plt.gca().invert_yaxis()
plt.xlabel("Importance (%)")
plt.ylabel("Features")
plt.title("Top 15 Most Important Features (Random Forest)")
plt.grid(True, axis='x', linestyle='--', alpha=0.7)
plt.show()
```

```
# Show the full table if needed  
display(feat_importance_df.head(20))
```

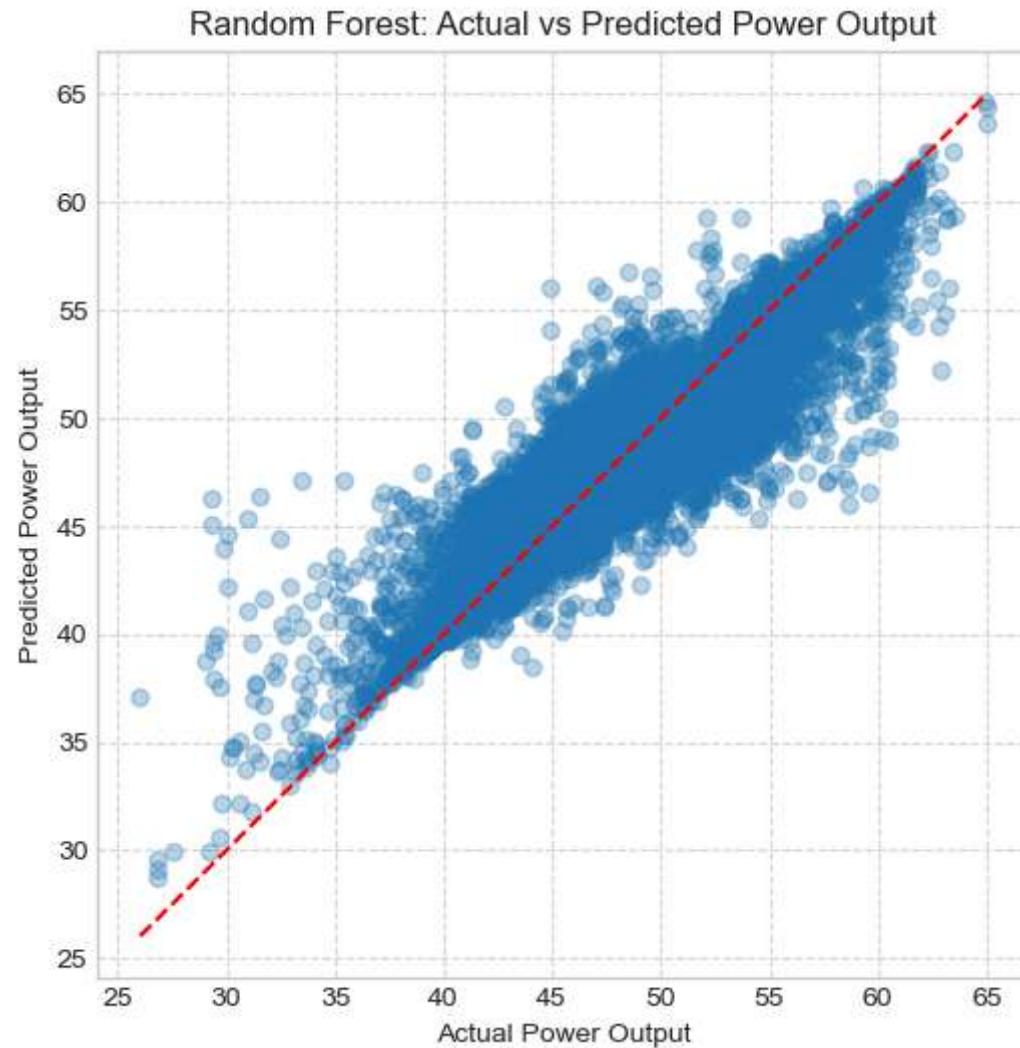


	Feature	Importance	Importance (%)
1	ambient_temperature	0.232586	23.258647
13	month	0.152748	15.274846
6	nacelle_temp	0.147612	14.761166
8	wind_direction_raw	0.087273	8.727287
3	generator_winding_temp_max	0.075651	7.565101
5	nc1_inside_temp	0.068514	6.851357
11	hour	0.054331	5.433093
14	season_-0.47581831258795315	0.037499	3.749900
4	grid_power10min_average	0.035436	3.543630
7	reactive_power	0.031291	3.129082
12	day_of_week	0.021431	2.143142
0	active_power_calculated_by_converter	0.017461	1.746082
10	wind_speed_turbulence	0.014302	1.430171
9	wind_speed_raw	0.011813	1.181262
2	generator_speed	0.008556	0.855624
15	season_0.4337021794943158	0.002661	0.266071
16	season_1.3432226715765847	0.000835	0.083541

In [71]:

```
# Actual vs Predicted plot
plt.figure(figsize=(6,6))
plt.scatter(y_test, y_pred_rf, alpha=0.3)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--')
plt.xlabel("Actual Power Output")
plt.ylabel("Predicted Power Output")
plt.title("Random Forest: Actual vs Predicted Power Output")
```

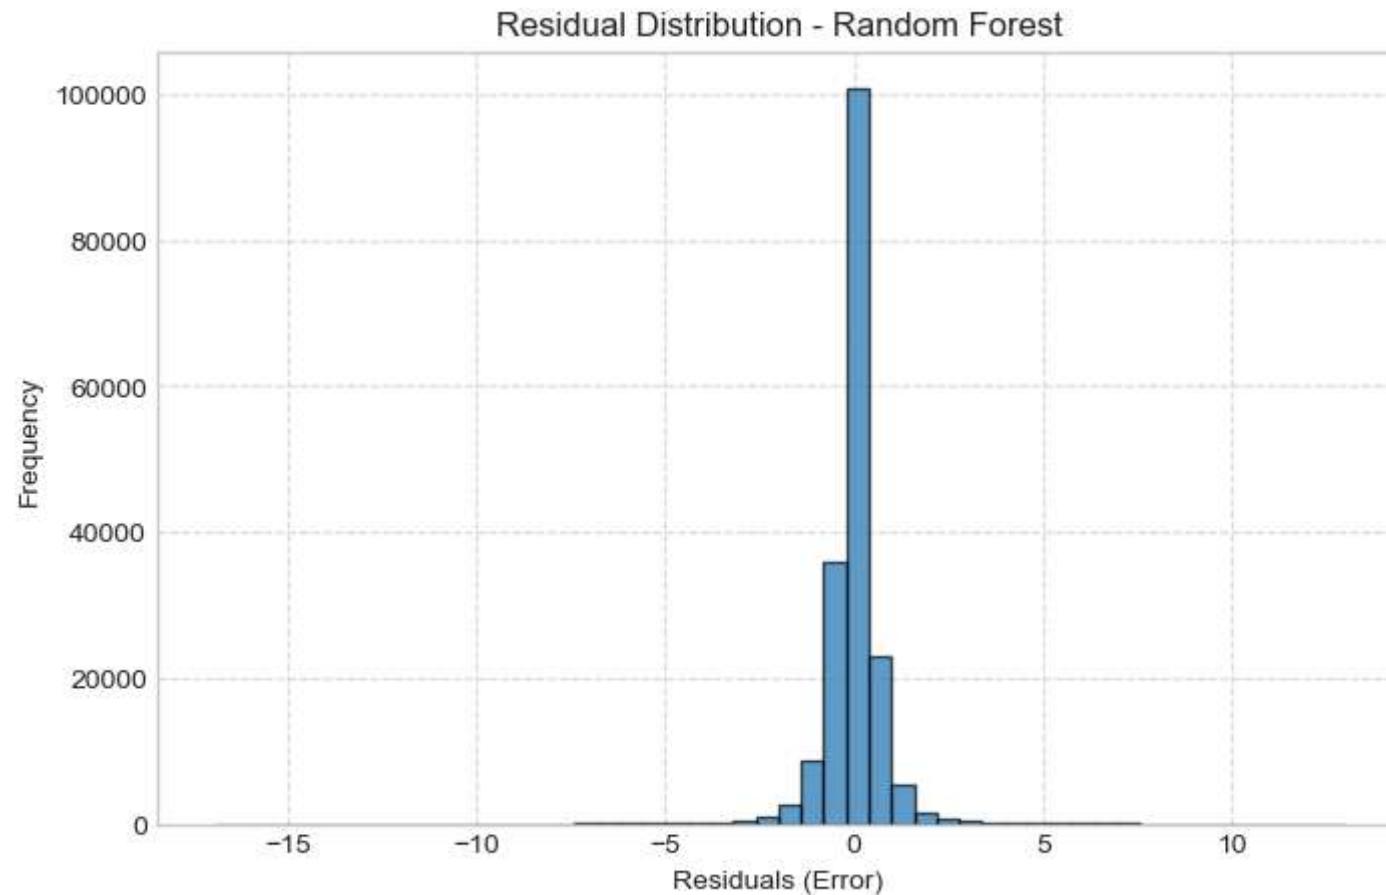
```
plt.grid(True, linestyle='--', alpha=0.7)
plt.show()
```



```
In [72]: # Residual Plot
residuals_rf = y_test - y_pred_rf

plt.figure(figsize=(8,5))
plt.hist(residuals_rf, bins=50, edgecolor='black', alpha=0.7)
```

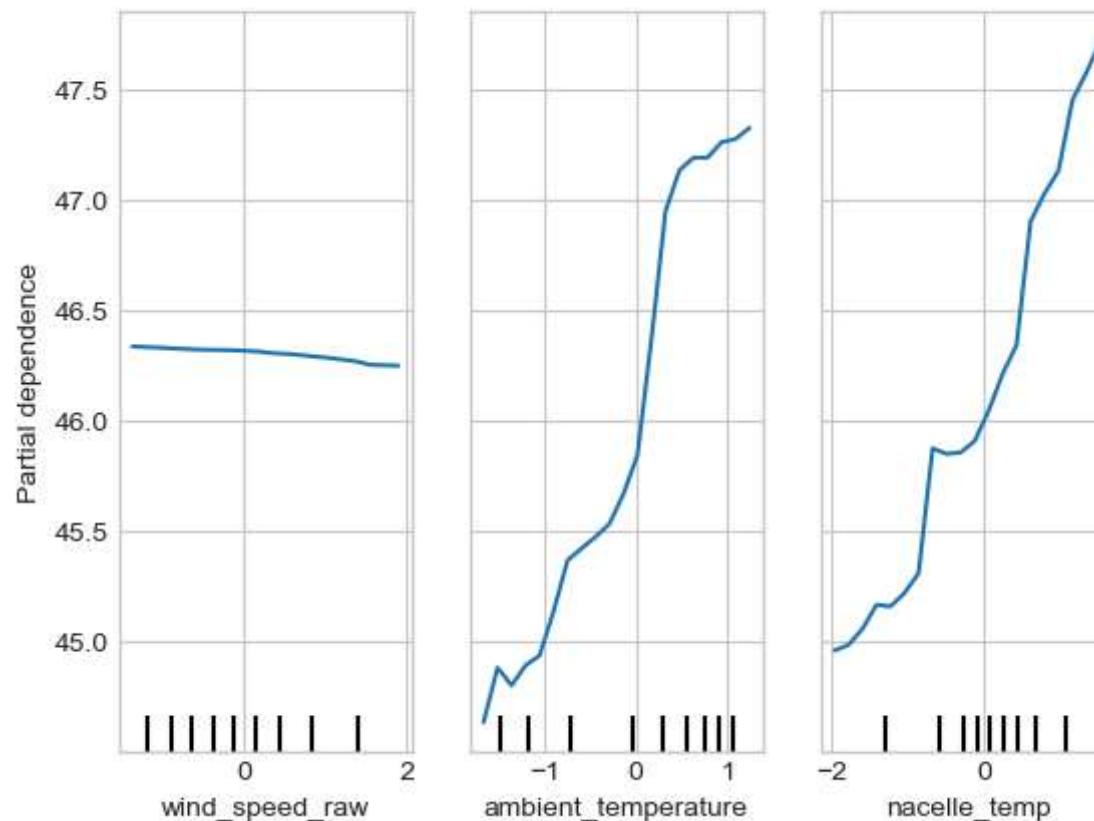
```
plt.xlabel("Residuals (Error)")
plt.ylabel("Frequency")
plt.title("Residual Distribution - Random Forest")
plt.grid(True, linestyle='--', alpha=0.6)
plt.show()
```



In [73]: # Partial Dependence Plots (PDPs)

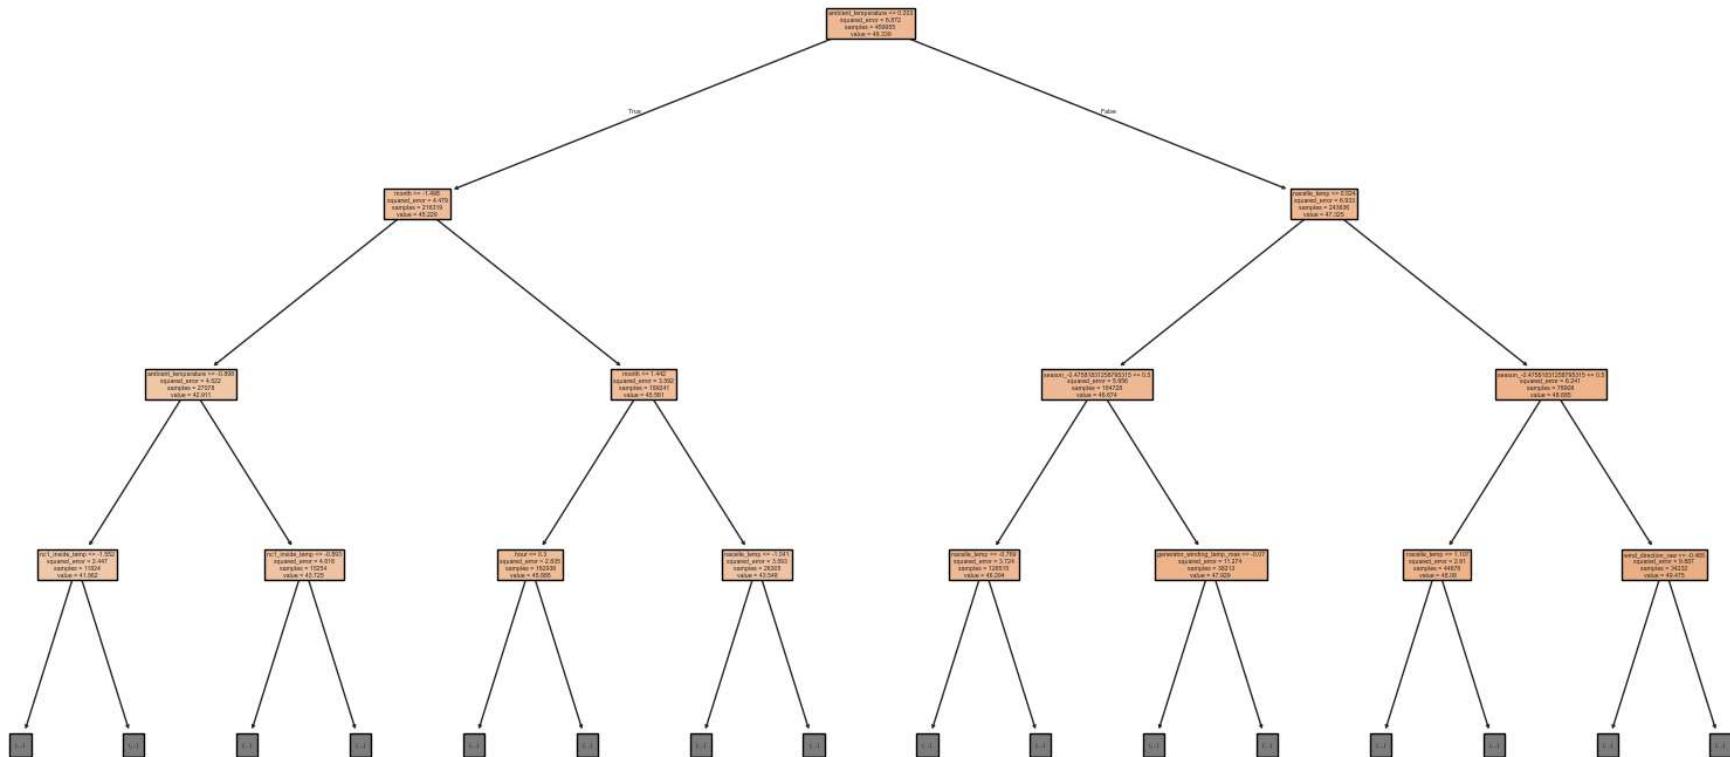
```
features_to_plot = ['wind_speed_raw', 'ambient_temperature', 'nacelle_temp']
PartialDependenceDisplay.from_estimator(rf_model, X_test, features_to_plot, grid_resolution=20)
plt.suptitle("Partial Dependence Plots (Random Forest)")
plt.show()
```

Partial Dependence Plots (Random Forest)



```
In [74]: # Tree Depth & Complexity  
  
plt.figure(figsize=(20,10))  
plot_tree(rf_model.estimators_[0], feature_names=X_train.columns, filled=True, max_depth=3)  
plt.title("Sample Decision Tree from Random Forest")  
plt.show()
```

Sample Decision Tree from Random Forest



```
In [75]: # Define tolerance ( within ±10% of actual value)
tolerance = 0.10
within_tolerance = np.abs((y_test - y_pred_rf) / y_test) <= tolerance
accuracy = within_tolerance.mean() * 100

print(f"Approximate Accuracy (within ±10% tolerance): {accuracy:.2f}%")
```

Approximate Accuracy (within ±10% tolerance): 99.71%

```
In [76]: # *****XGBoost Regressor*****
```

```
In [77]: # Train XGBoost Regressor
xgb_model = XGBRegressor(
    n_estimators=300,
    learning_rate=0.05,
    max_depth=8,
    subsample=0.8,
    colsample_bytree=0.8,
    random_state=42,
    n_jobs=-1
)
```

```
In [78]: xgb_model.fit(X_train, y_train)
```

```
Out[78]:
```

XGBRegressor

```
XGBRegressor(base_score=None, booster=None, callbacks=None,
             colsample_bylevel=None, colsample_bynode=None,
             colsample_bytree=0.8, device=None, early_stopping_rounds=None,
             enable_categorical=False, eval_metric=None, feature_types=None,
             feature_weights=None, gamma=None, grow_policy=None,
             importance_type=None, interaction_constraints=None,
             learning_rate=0.05, max_bin=None, max_cat_threshold=None,
             max_cat_to_onehot=None, max_delta_step=None, max_depth=8,
             max_leaves=None, min_child_weight=None, missing=nan,
```

```
In [79]: # Predictions
y_pred_xgb = xgb_model.predict(X_test)
```

```
In [80]: # Metrics
mae_xgb = mean_absolute_error(y_test, y_pred_xgb)
rmse_xgb = np.sqrt(mean_squared_error(y_test, y_pred_xgb))
r2_xgb = r2_score(y_test, y_pred_xgb)

n = len(y_test)
p = X_test.shape[1]
adj_r2_xgb = 1 - ((1 - r2_xgb) * (n - 1) / (n - p - 1))
```

```
In [81]: # ±10% Tolerance Accuracy
tolerance = 0.10
within_tolerance = np.abs((y_test - y_pred_xgb) / y_test) <= tolerance
accuracy_xgb = within_tolerance.mean() * 100
```

```
In [82]: # Results
print("XGBoost Model Performance")
print("-----")
print(f"MAE : {mae_xgb:.4f}")
print(f"RMSE : {rmse_xgb:.4f}")
print(f"R² : {r2_xgb:.4f}")
print(f"Adj R²: {adj_r2_xgb:.4f}")
print(f"Approx Accuracy (±10% tolerance): {accuracy_xgb:.2f}%")
```

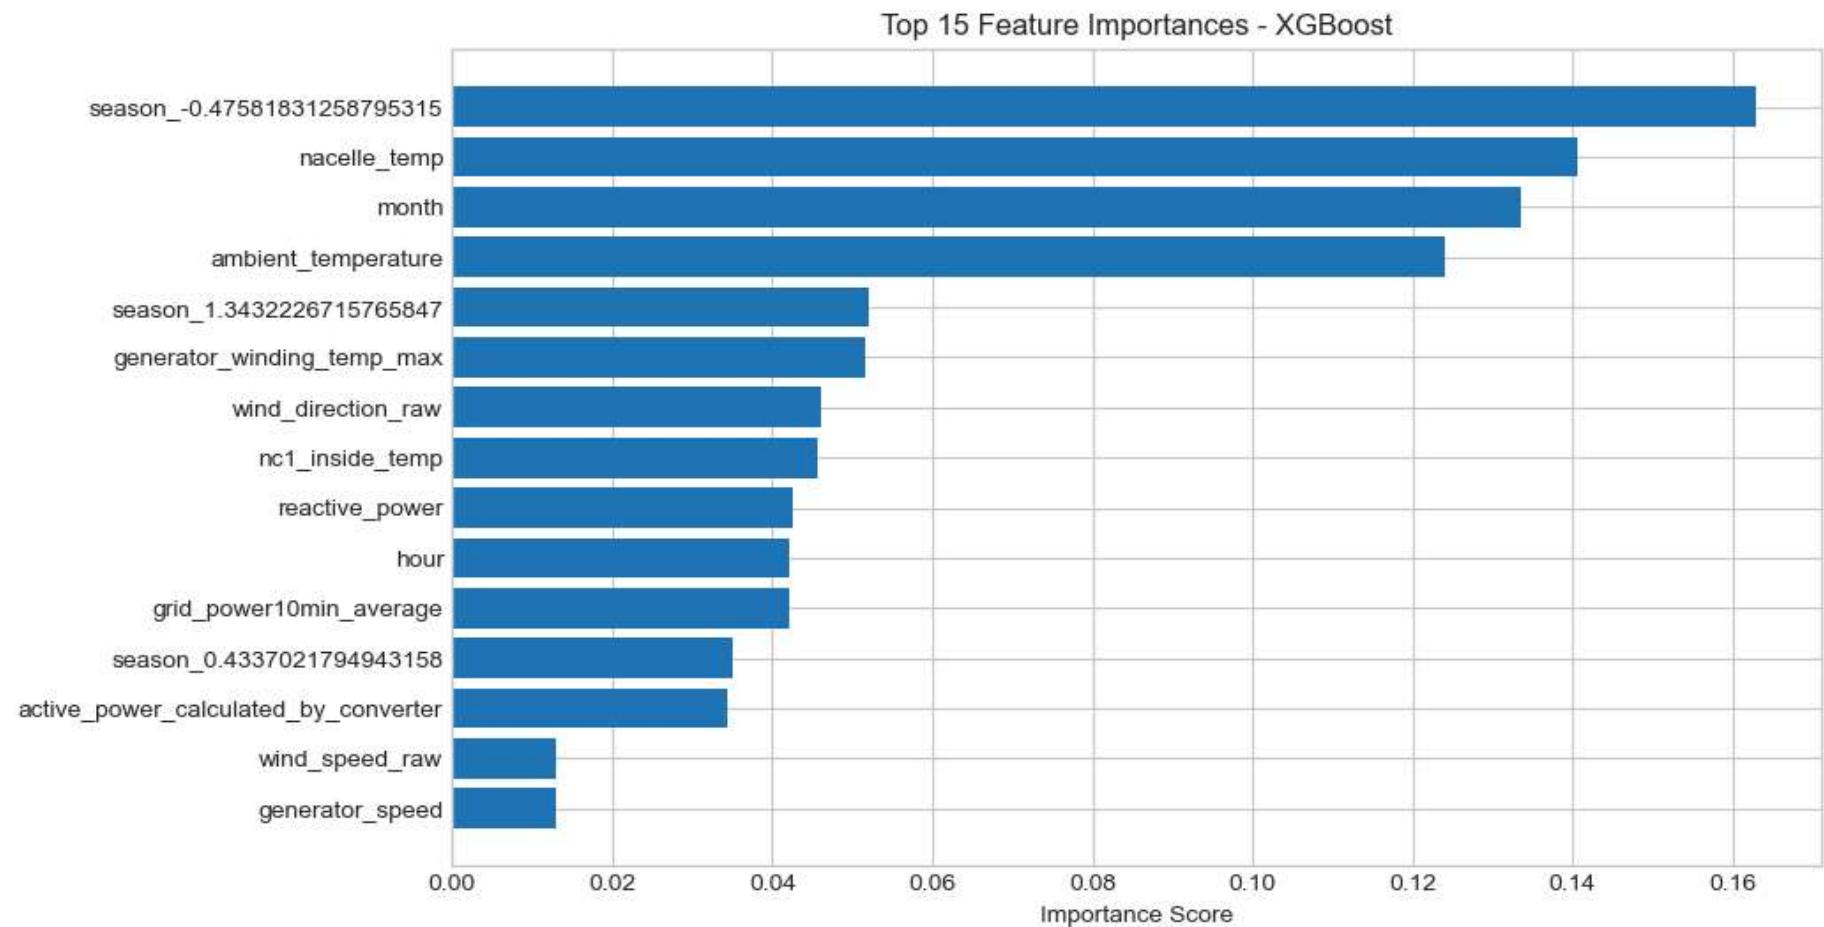
XGBoost Model Performance

MAE : 0.7747
RMSE : 1.1339
R² : 0.8117
Adj R²: 0.8117
Approx Accuracy (±10% tolerance): 99.44%

```
In [83]: # Feature Importance Visualization
importances = xgb_model.feature_importances_
feature_importance_xgb = pd.DataFrame({
    'Feature': X_train.columns,
    'Importance': importances
}).sort_values(by='Importance', ascending=False)
```

```
In [84]: plt.figure(figsize=(10,6))
plt.barh(feature_importance_xgb[ 'Feature' ].head(15), feature_importance_xgb[ 'Importance' ].head(15))
plt.gca().invert_yaxis()
plt.title("Top 15 Feature Importances - XGBoost")
plt.xlabel("Importance Score")
plt.show()

display(feature_importance_xgb.head(20))
```

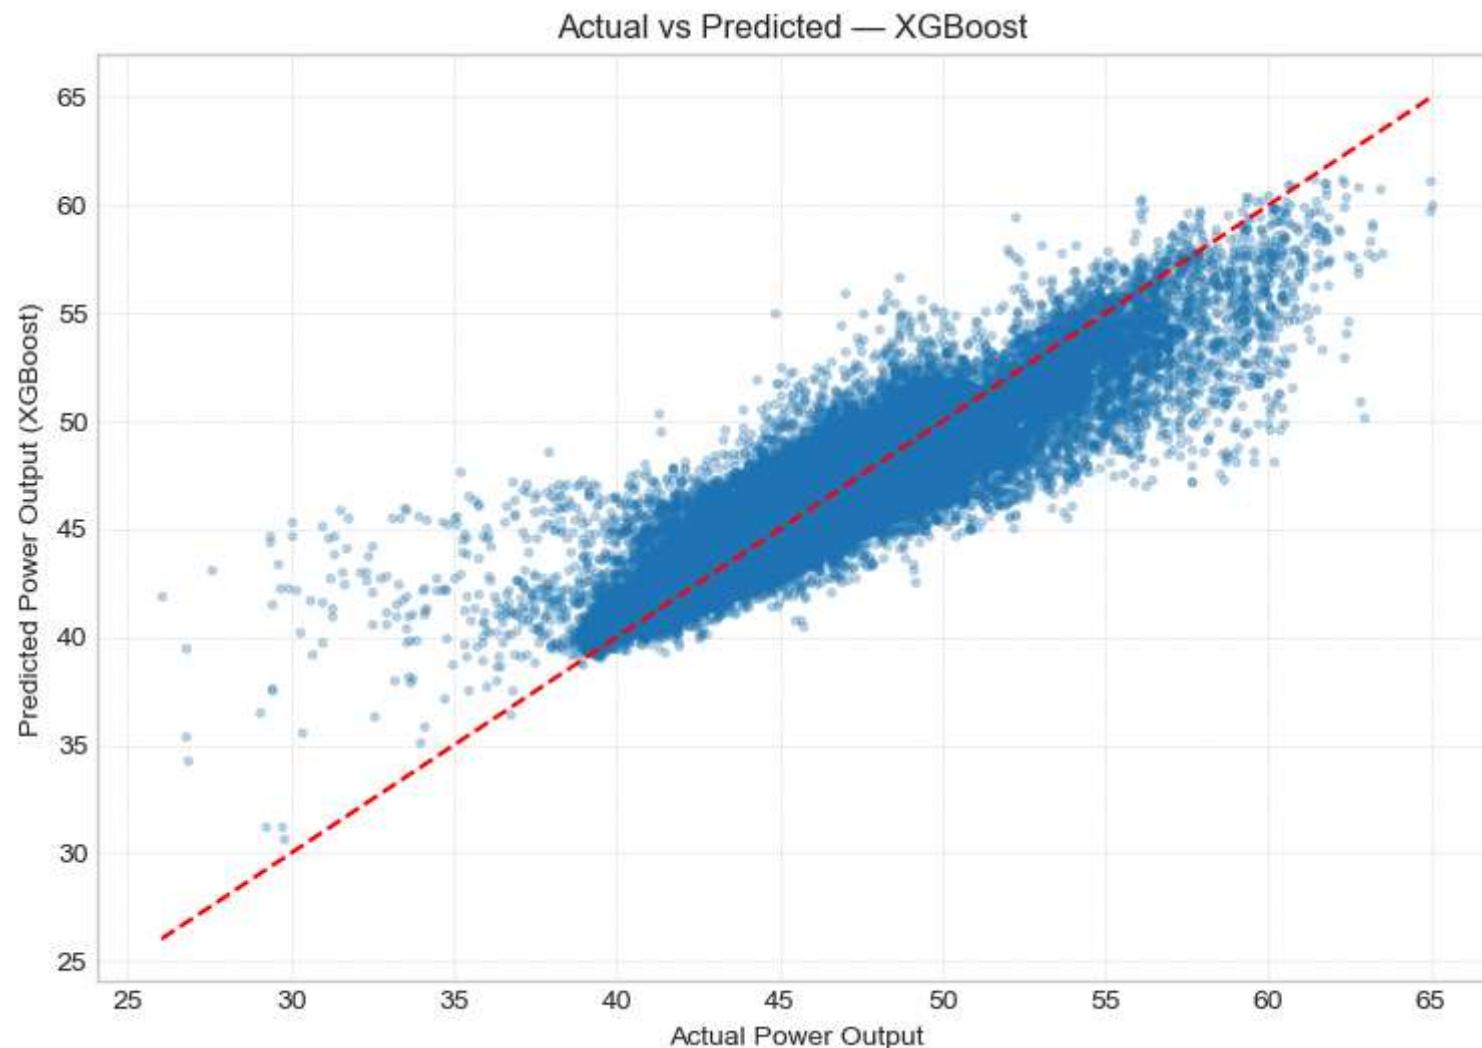


	Feature	Importance
14	season_-0.47581831258795315	0.163026
6	nacelle_temp	0.140588
13	month	0.133498
1	ambient_temperature	0.123960
16	season_1.3432226715765847	0.052126
3	generator_winding_temp_max	0.051552
8	wind_direction_raw	0.046034
5	nc1_inside_temp	0.045688
7	reactive_power	0.042583
11	hour	0.042103
4	grid_power10min_average	0.042040
15	season_0.4337021794943158	0.035077
0	active_power_calculated_by_converter	0.034444
9	wind_speed_raw	0.012990
2	generator_speed	0.012873
12	day_of_week	0.012290
10	wind_speed_turbulence	0.009126

```
In [85]: # ----- Actual vs Predicted -----
y_pred_series = pd.Series(y_pred_xgb, index=y_test.index)

plt.figure(figsize=(9,6))
plt.scatter(y_test, y_pred_series, alpha=0.25, s=7)
minval = min(y_test.min(), y_pred_series.min())
maxval = max(y_test.max(), y_pred_series.max())
```

```
plt.plot([minval, maxval], [minval, maxval], 'r--')
plt.xlabel("Actual Power Output")
plt.ylabel("Predicted Power Output (XGBoost)")
plt.title("Actual vs Predicted – XGBoost ")
plt.grid(True, linestyle=':', linewidth=0.5)
plt.show()
```



In [86]: `residuals_ = y_test - y_pred_series`

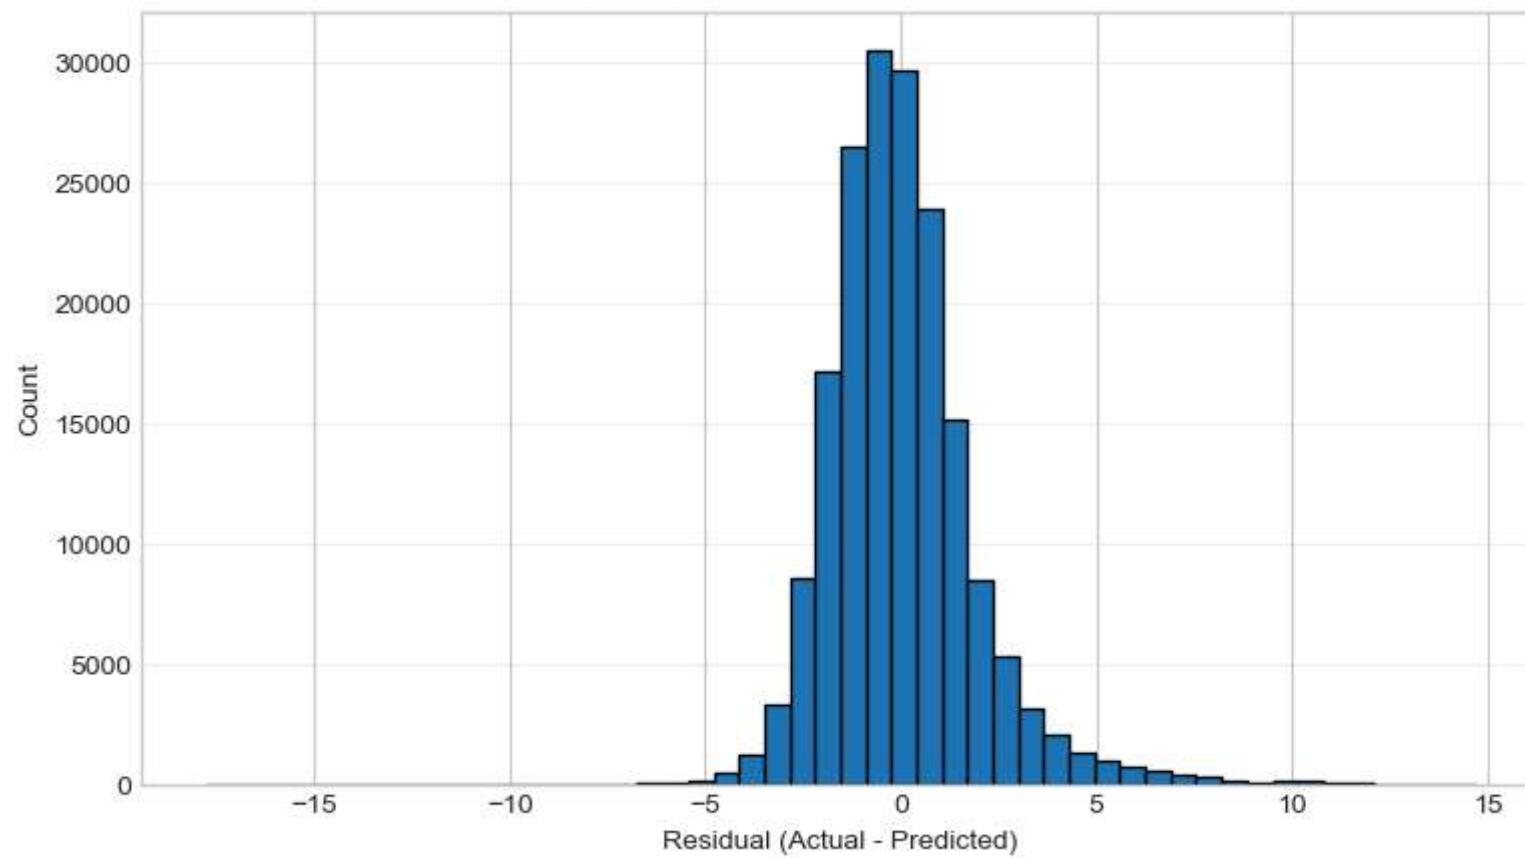
```
In [87]: # Quick check to print a few large errors
err_abs = np.abs(residuals)
top_err = err_abs.sort_values(ascending=False).head(10)
print("Top 10 absolute errors (index -> error):")
print(top_err)
```

```
Top 10 absolute errors (index -> error):
531587    17.786339
822445    17.508259
20052     16.708145
571686    16.486479
150998    16.366166
18711     16.321634
104783    16.279990
70554     16.231050
84697     16.101582
904449    15.986548
dtype: float64
```

```
In [88]: # Residuals histogram
plt.figure(figsize=(9,5))
plt.hist(residuals, bins=50, edgecolor='black')
plt.xlabel("Residual (Actual - Predicted)")
plt.ylabel("Count")
plt.title("Residual Distribution - XGBoost")
plt.grid(axis='y', linestyle=':', linewidth=0.5)
plt.show()

print("Residuals mean: {:.4f}, median: {:.4f}, std: {:.4f}".format(residuals.mean(), residuals.median(), residuals.std()))
```

Residual Distribution — XGBoost

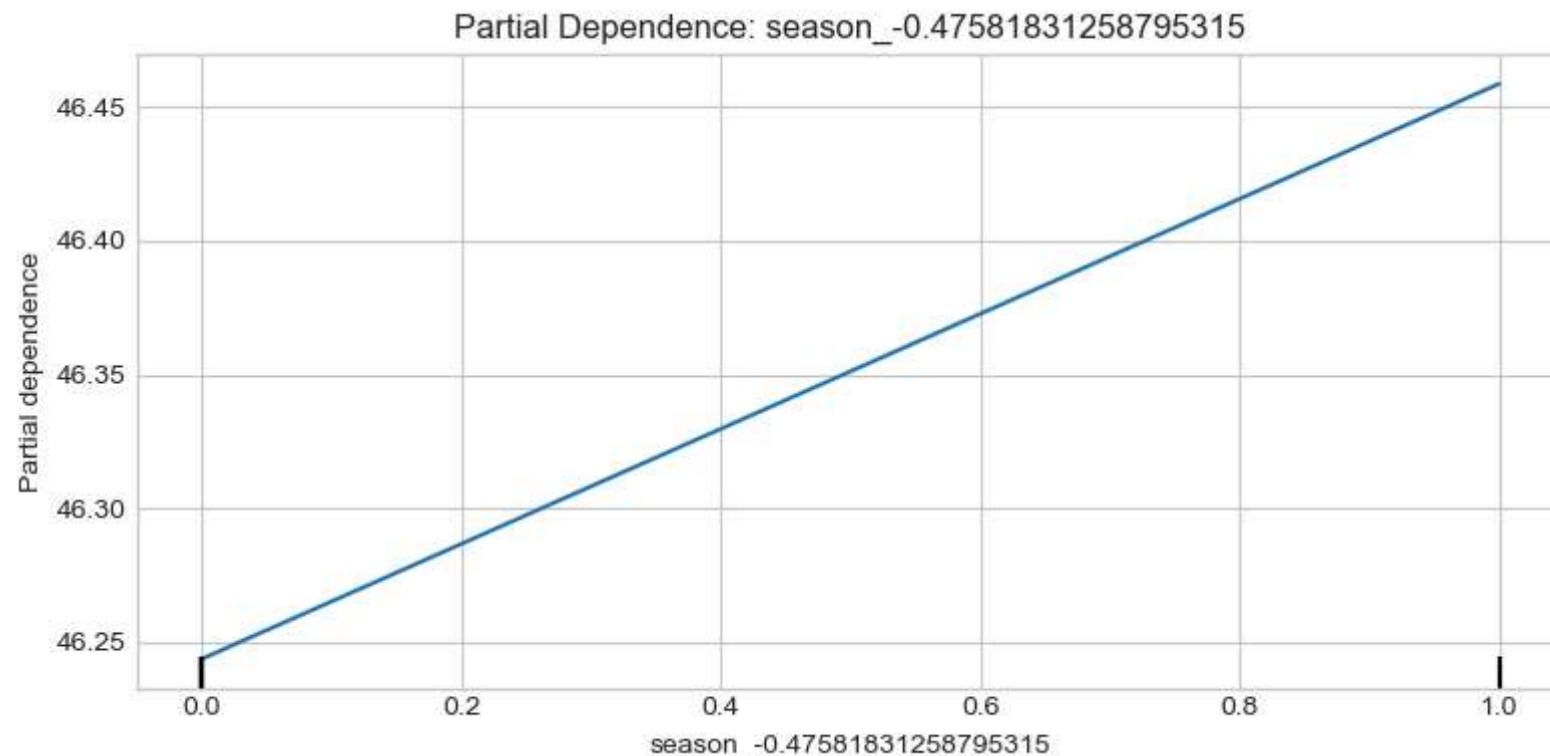


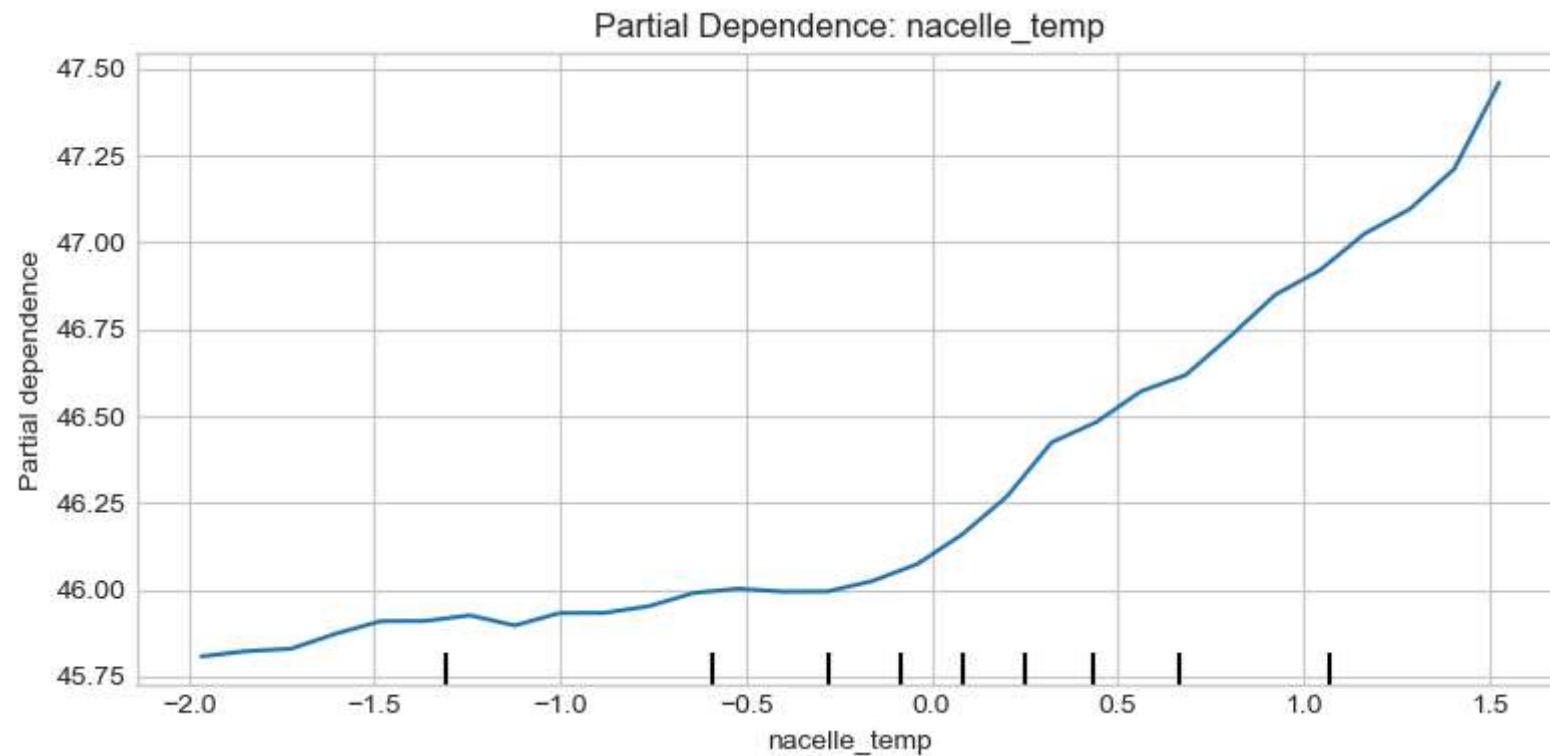
Residuals mean: -0.0052, median: -0.1984, std: 1.9361

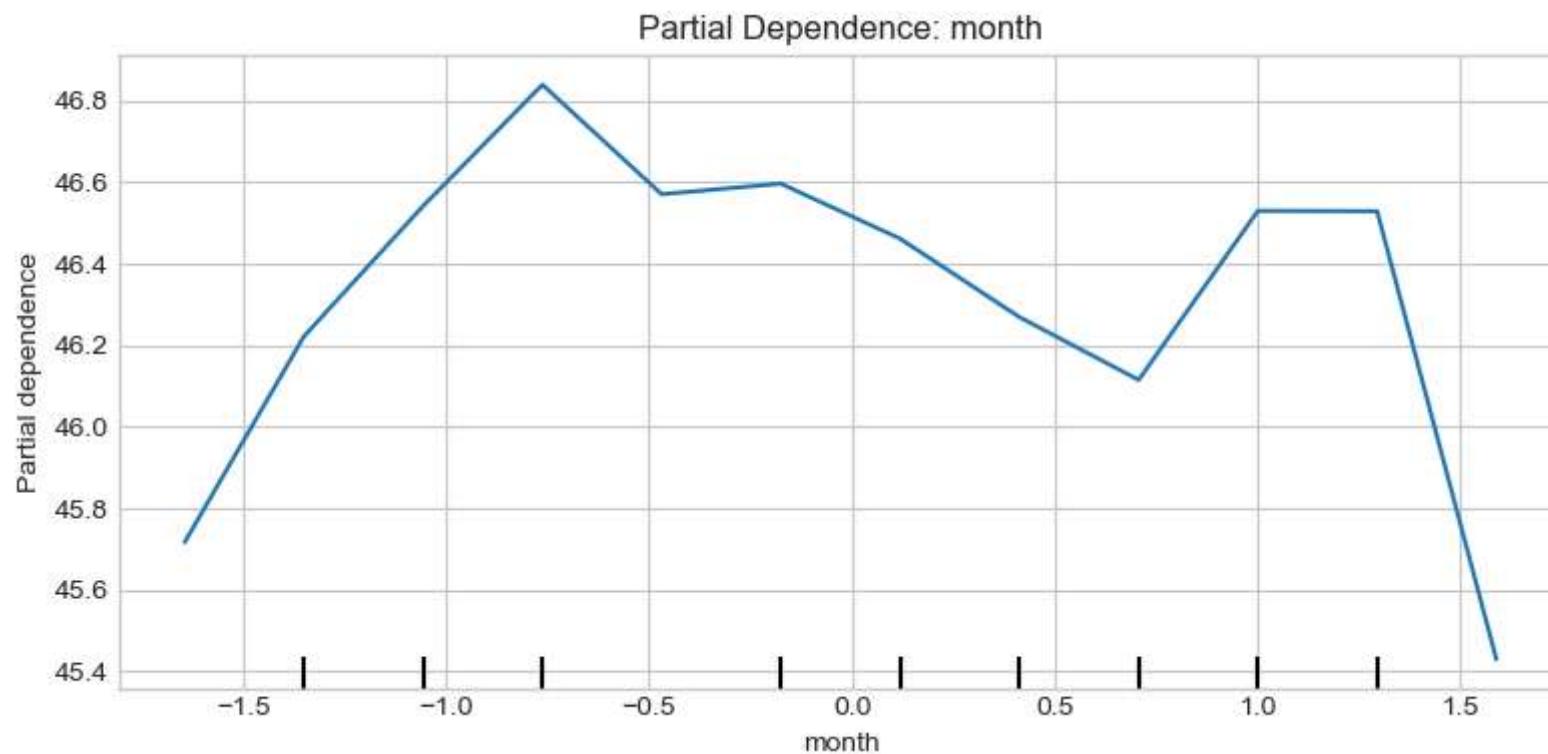
```
In [89]: # Partial Dependence Plots (top 3 features)
# Filter to ensure features exist in X_test
top3 = feature_importance_xgb['Feature'].head(3).tolist()
top3 = [f for f in top3 if f in X_test.columns]
if len(top3) == 0:
    raise ValueError("No valid PDP features found in X_test. Check `feature_importance` or X_test columns.")
```

```
In [90]: for feat in top3:
    plt.figure(figsize=(8,4))
    PartialDependenceDisplay.from_estimator(xgb_model, X_test, [feat], grid_resolution=30, ax=plt.gca())
    plt.title(f"Partial Dependence: {feat}")
```

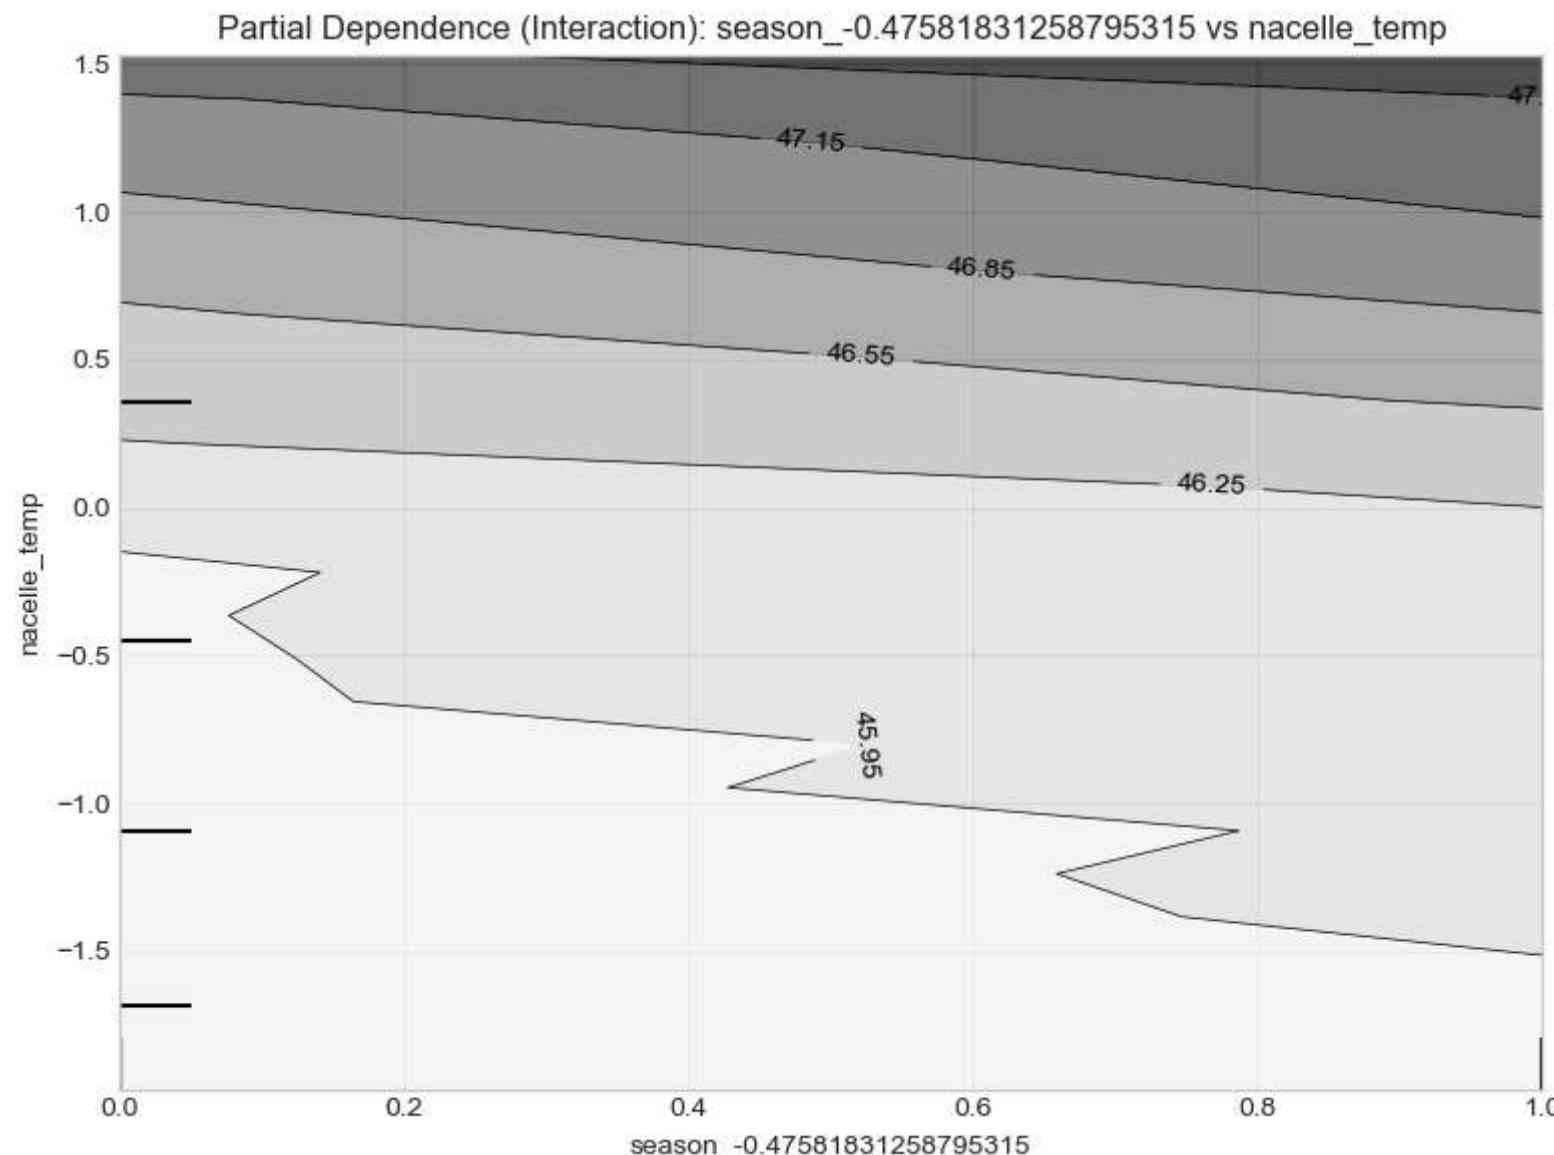
```
plt.tight_layout()  
plt.show()
```







```
In [91]: # Interaction PDP between top1 & top2
if len(top3) >= 2:
    pair = (top3[0], top3[1])
    plt.figure(figsize=(8,6))
    PartialDependenceDisplay.from_estimator(xgb_model, X_test, [pair], grid_resolution=25, ax=plt.gca())
    plt.title(f"Partial Dependence (Interaction): {pair[0]} vs {pair[1]}")
    plt.tight_layout()
    plt.show()
```



```
In [92]: # Model Comparision
```

```
In [93]: # ----- Checking if the models are overfitting or not-----
```

```
# Predict on training and test data
y_train_pred_lr = lr_model.predict(X_train_scaled)
y_test_pred_lr = lr_model.predict(X_test_scaled)

y_train_pred_rf = rf_model.predict(X_train_scaled)
y_test_pred_rf = rf_model.predict(X_test_scaled)

y_train_pred_xgb = xgb_model.predict(X_train_scaled)
y_test_pred_xgb = xgb_model.predict(X_test_scaled)

# -----Compute metrics-----

# -----Training metrics-----

# Linear Regression
mae_train_lr = mean_absolute_error(y_train, y_train_pred_lr)
rmse_train_lr = np.sqrt(mean_squared_error(y_train, y_train_pred_lr))
r2_train_lr = r2_score(y_train, y_train_pred_lr)

# Random Forest Regressor
mae_train_rf = mean_absolute_error(y_train, y_train_pred_rf)
rmse_train_rf = np.sqrt(mean_squared_error(y_train, y_train_pred_rf))
r2_train_rf = r2_score(y_train, y_train_pred_rf)

# XGBoost
mae_train_xgb = mean_absolute_error(y_train, y_train_pred_xgb)
rmse_train_xgb = np.sqrt(mean_squared_error(y_train, y_train_pred_xgb))
r2_train_xgb = r2_score(y_train, y_train_pred_xgb)

# -----Testing metrics-----

# Linear Regression
mae_test_lr = mean_absolute_error(y_test, y_test_pred_lr)
rmse_test_lr = np.sqrt(mean_squared_error(y_test, y_test_pred_lr))
r2_test_lr = r2_score(y_test, y_test_pred_lr)

# Random Forest Regressor
mae_test_rf = mean_absolute_error(y_test, y_test_pred_rf)
rmse_test_rf = np.sqrt(mean_squared_error(y_test, y_test_pred_rf))
r2_test_rf = r2_score(y_test, y_test_pred_rf)
```

```

# XGBoost
mae_test_xgb = mean_absolute_error(y_test, y_test_pred_xgb)
rmse_test_xgb = np.sqrt(mean_squared_error(y_test, y_test_pred_xgb))
r2_test_xgb = r2_score(y_test, y_test_pred_xgb)

# -----Linear Regression-----
print("Linear regression Performance Comparison")
print("-----")
print(f"Train MAE : {mae_train_lr:.4f} | Test MAE : {mae_test_lr:.4f}")
print(f"Train RMSE : {rmse_train_lr:.4f} | Test RMSE : {rmse_test_lr:.4f}")
print(f"Train R² : {r2_train_lr:.4f} | Test R² : {r2_test_lr:.4f}")

tolerance = 0.10
acc_train_lr = (np.abs((y_train - y_train_pred_lr) / y_train) <= tolerance).mean() * 100
acc_test_lr = (np.abs((y_test - y_test_pred_lr) / y_test) <= tolerance).mean() * 100
print(f"Train Accuracy (\u00b110%): {acc_train_lr:.2f}% | Test Accuracy (\u00b110%): {acc_test_lr:.2f%}\n\n")

# -----Random Forest Regressor-----
print("Random Forest Performance Comparison")
print("-----")
print(f"Train MAE : {mae_train_rf:.4f} | Test MAE : {mae_test_rf:.4f}")
print(f"Train RMSE : {rmse_train_rf:.4f} | Test RMSE : {rmse_test_rf:.4f}")
print(f"Train R² : {r2_train_rf:.4f} | Test R² : {r2_test_rf:.4f}")

acc_train_rf = (np.abs((y_train - y_train_pred_rf) / y_train) <= tolerance).mean() * 100
acc_test_rf = (np.abs((y_test - y_test_pred_rf) / y_test) <= tolerance).mean() * 100
print(f"Train Accuracy (\u00b110%): {acc_train_rf:.2f}% | Test Accuracy (\u00b110%): {acc_test_rf:.2f%} \n\n")

# -----XGBoost-----
print("XGBoost Performance Comparison")
print("-----")
print(f"Train MAE : {mae_train_xgb:.4f} | Test MAE : {mae_test_xgb:.4f}")
print(f"Train RMSE : {rmse_train_xgb:.4f} | Test RMSE : {rmse_test_xgb:.4f}")
print(f"Train R² : {r2_train_xgb:.4f} | Test R² : {r2_test_xgb:.4f}")

acc_train_xgb = (np.abs((y_train - y_train_pred_xgb) / y_train) <= tolerance).mean() * 100

```

```
acc_test_xgb = (np.abs((y_test - y_test_pred_xgb) / y_test) <= tolerance).mean() * 100
print(f"Train Accuracy (\u00b110%): {acc_train_xgb:.2f}% | Test Accuracy (\u00b110%): {acc_test_xgb:.2f}%\n\n")
```

Linear regression Performance Comparison

```
-----
Train MAE : 1.3848 | Test MAE : 1.3793
Train RMSE : 1.9452 | Test RMSE : 1.9361
Train R2 : 0.4483 | Test R2 : 0.4511
Train Accuracy (\u00b110%): 97.44% | Test Accuracy (\u00b110%): 97.48%
```

Random Forest Performance Comparison

```
-----
Train MAE : 0.1625 | Test MAE : 0.4329
Train RMSE : 0.2901 | Test RMSE : 0.7609
Train R2 : 0.9877 | Test R2 : 0.9152
Train Accuracy (\u00b110%): 99.99% | Test Accuracy (\u00b110%): 99.71%
```

XGBoost Performance Comparison

```
-----
Train MAE : 0.7536 | Test MAE : 0.7747
Train RMSE : 1.0949 | Test RMSE : 1.1339
Train R2 : 0.8252 | Test R2 : 0.8117
Train Accuracy (\u00b110%): 99.51% | Test Accuracy (\u00b110%): 99.44%
```

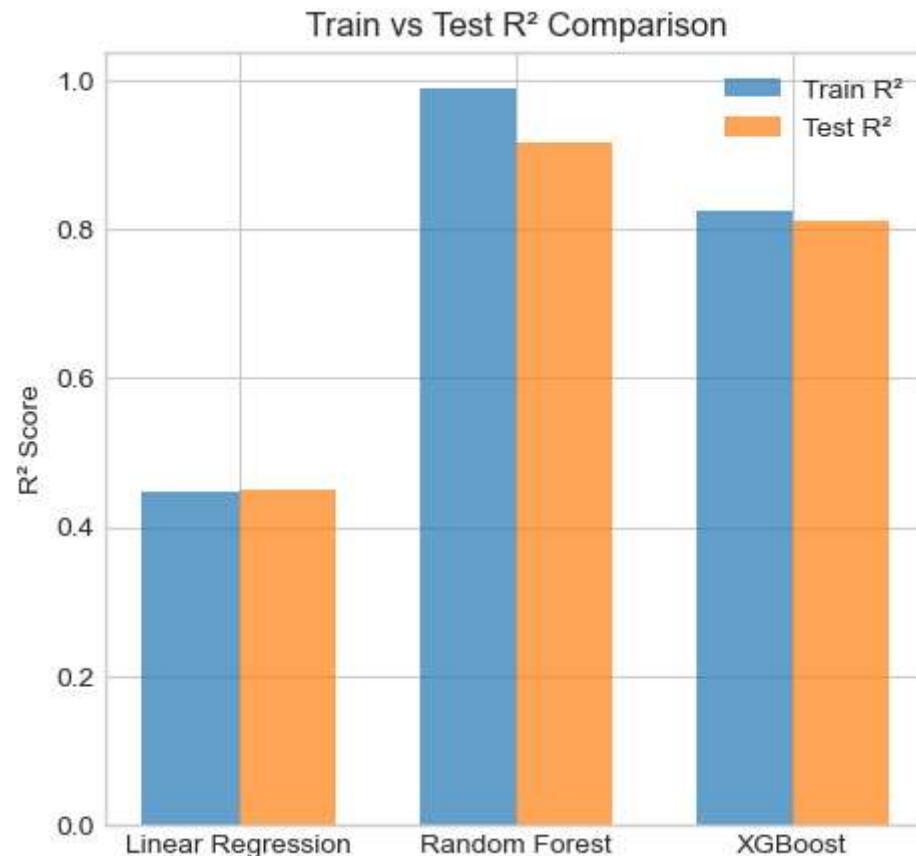
```
In [94]: # ----- PLOT 1: R2 Comparison -----
plt.figure(figsize=(18, 5))

plt.subplot(1, 3, 1)
models = ['Linear Regression', 'Random Forest', 'XGBoost']
train_scores = [r2_train_lr, r2_train_rf, r2_train_xgb]
test_scores = [r2_test_lr, r2_test_rf, r2_test_xgb]

x = np.arange(len(models))
width = 0.35

plt.bar(x - width/2, train_scores, width, label='Train R2', alpha=0.7)
plt.bar(x + width/2, test_scores, width, label='Test R2', alpha=0.7)
```

```
plt.title("Train vs Test R2 Comparison")
plt.ylabel("R2 Score")
plt.xticks(x, models)
plt.legend()
plt.show()
```



```
In [95]: # ----- Accuracy VS Tolerance Curve-----
preds = {
    "Linear Regression": y_pred_lr,
    "Random Forest": y_pred_rf,
    "XGBoost": y_pred_xgb
}

# Actual values
```

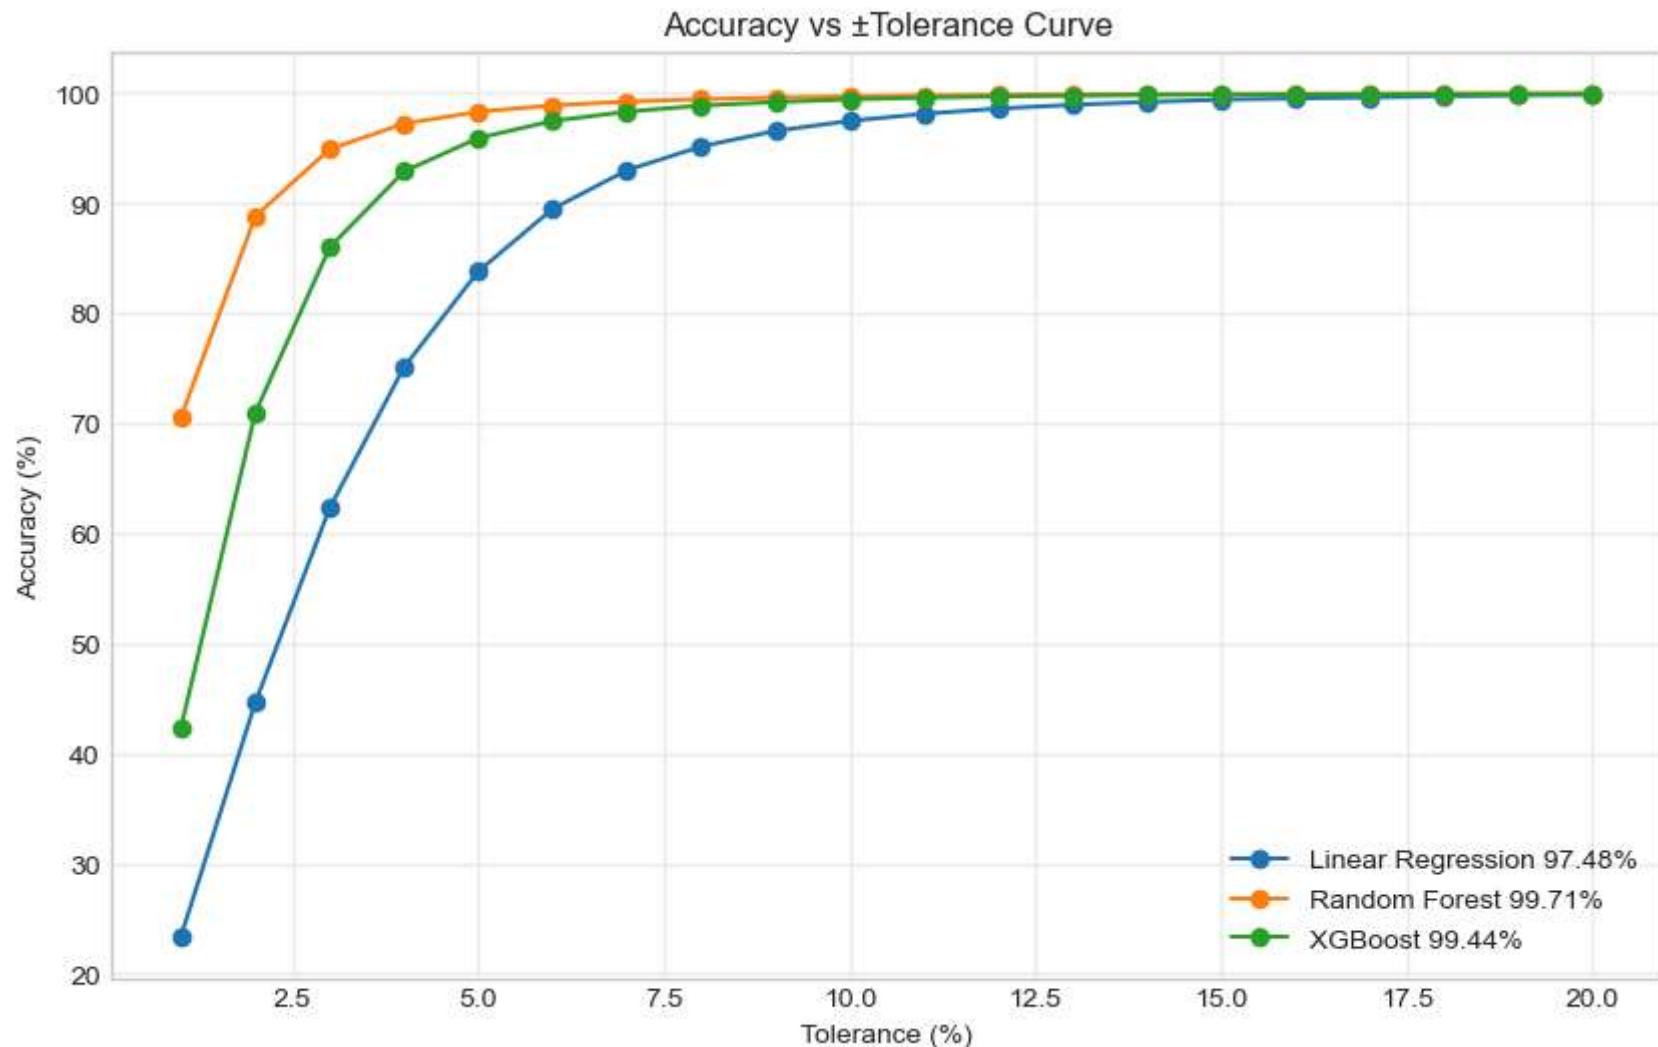
```
actual = y_test.values

# Define tolerance range (1% to 20)
tolerances = np.linspace(0.01, 0.20, 20)

# Compute accuracy within each tolerance for each model
accuracies = {}
for name, y_pred in preds.items():
    model_acc = []
    for tol in tolerances:
        within = np.abs((actual - y_pred) / actual) <= tol
        model_acc.append(within.mean() * 100)
    accuracies[name] = model_acc

# Plot
plt.figure(figsize=(10,6))
for name, acc in accuracies.items():
    idx = np.argmin(np.abs(tolerances - 0.10))
    plt.plot(tolerances * 100, acc, 'o-', label=f'{name} {acc[idx]:.2f}%')

plt.title("Accuracy vs ±Tolerance Curve")
plt.xlabel("Tolerance (%)")
plt.ylabel("Accuracy (%)")
plt.legend()
plt.grid(alpha=0.3)
plt.show()
```



```
In [96]: # ----- Feature Importance Comparison -----
lr_coef_df_sorted = coef_df.head(10)

rf_importances_sorted = feat_importance_df.head(10)

xgb_importances_sorted = feature_importance_xgb.head(10)

plt.figure(figsize = (20, 6))
```

```
plt.subplot(1, 3, 3)

dataframes_to_merge = [
    lr_coef_df_sorted[['Feature', 'Coefficient']],
    rf_importances_sorted[['Feature', 'Importance']],
    xgb_importances_sorted[['Feature', 'Importance']]
]

top_features = reduce(
    lambda left, right: pd.merge(left, right, on='Feature', how='outer'),
    dataframes_to_merge
).fillna(0).head(10)

top_features.plot(
    x='Feature',
    kind='bar',
    ax=plt.gca(),
    color=['steelblue', 'orange', 'green'],
    alpha=0.8
)

plt.title("Top Feature Comparison (LR Coeff vs RF Importance VS XBG Importance)")
plt.xticks(rotation=45, ha='right')
plt.ylabel("Weight / Importance")
plt.legend(['LR Coefficient', 'RF Importance', 'XGB Importance'])

plt.tight_layout()
plt.show()
```

